



STW7071CEM

INFORMATION RETRIEVAL



-  Search
-  Retrieve
-  Cluster
-  Relevant Results



STUDENT NAME | Rohit Jha



COVENTRY ID | 11782276



SOFTWARICA ID | 210178



MODULE NAME | Information Retrieval



MODULE CODE | ST7071CEM



MODULE LEADER | Siddhartha Neupane



Table of Contents

List of Abbreviations	4
Conceptual System Diagram.....	5
Introduction.....	6
Task 1 - Vertical Search Engine	7
2.1 Introduction and Problem Statement	7
2.2 Objectives	7
2.3 System Architecture.....	7
2.4 Web Crawler Component	7
Robots.txt Compliance and Ethical Crawling.....	8
2.5 Scheduling and Automation.....	8
2.6 Data Collection and MongoDB Storage	9
2.7 Text Preprocessing Pipeline.....	10
2.8 Vector Space Model.....	11
2.9 Cosine Similarity and Ranking	12
2.10 Query Processing	13
2.11 Top-K Search and Pagination	13
2.12 Search Interface	13
2.13 Evaluation and Testing	15
Search Retrieval Evaluation.....	15
2.14 Limitations and Improvements	16
Task 2 - News Document Clustering	17
3.1 Introduction and Problem Statement	17
3.2 Objectives	17
3.3 Literature Review and Theoretical Foundation	17
Principal Component Analysis (PCA).....	17
3.4 Dataset and Data Engineering.....	18
3.5 Text Preprocessing and Feature Engineering	18
3.6 TF-IDF Vectorisation.....	18
3.7 K-Means Clustering Methodology	19
3.8 User Input Classification.....	20
3.9 MongoDB Storage	21
3.10 Clustering Visualisation.....	22

3.11 Results and Performance Analysis.....	24
3.12 Confusion Matrix & Heatmap Evaluation	25
3.13 Discussion and Findings	26
Overall Discussion	26
Conclusion	29
References.....	30
Appendix.....	31
Appendix A: Project Links	31
Appendix B: Architecture & File Structure	31
Appendix C: Full Code Screenshots	32
Appendix D: Pytest Evidence	89
Appendix E: Additional Application Screenshots	90

List of Figures & Tables

Figure 1 Unified System Architecture for Search and Clustering	5
Figure 2 Code snippet showing robots.txt parsing and the 5-second Crawl-Delay enforcement.....	8
Figure 3 Code snippet for polite HTTP retrieval ensuring crawler etiquette.	9
Figure 4 MongoDB Compass screenshot showing the research_outputs collection populated with crawled documents.....	10
Figure 5 Demonstrating tokenisation, lowercase folding, and stop-word removal for the crawler.....	11
Figure 6 Python code snippet demonstrating the TF-IDF vector matrix construction and storage.	12
Figure 7 Python code snippet showing the cosine similarity dot product operation. ..	13
Figure 8: OneSpot.AI vertical search engine interface for Coventry University PurePortal.....	13
Figure 9 Live crawler status UI showing indexed documents and real-time backend updates.	14
Figure 10 Auto-suggestion functionality for researcher names.	14
Figure 11 Search results for 'mental health' ranked by cosine similarity.....	15

Figure 12 Python code snippet from build_dataset.py explicitly loading the approved academic BBC dataset.	18
Figure 13 Code snippet demonstrating Text Preprocessing applied to the corpus.	18
Figure 14 Elbow Method - within-cluster SSE (inertia) vs. number of clusters K. The inflection point at K=3 empirically validates the choice of three clusters.....	19
Figure 15 Code snippet demonstrating the K-Means model initialisation and training pipeline.....	20
Figure 16 Code snippet from the backend API demonstrating live K-Means prediction on user input.....	21
Figure 17 MongoDB Compass screenshot showing the clustering_predictions collection.....	22
Figure 18 Document classification GUI providing live predictions.....	22
Figure 19 Help Modal containing pre-engineered test phrases per category.....	23
Figure 20 Prediction History log displaying user inputs and classifications.	23
Figure 21 Demonstration of Data Bias: 'President' classified as Economics.....	23
Figure 22 PCA 2D projection of the K-Means document clusters, illustrating the three distinct topic boundaries.	25
Figure 23 Heatmap visualisation of the Confusion Matrix. The diagonal proves high predictive accuracy, while the off-diagonal cells reveal the natural semantic overlap between Politics and Economics.....	26
Figure 24 Complete project file and folder structure.....	31
Figure 25 Full Pytest execution for task 1	89
Figure 26 Full Pytest execution for Task 2.....	89

Table 1 Summary of MongoDB collections and document counts after a successful crawl.....	9
Table 2 MongoDB Schema definition for the research_outputs collection.....	10
Table 3 Empirical PyTest unit-testing results for the Task 1 crawler and ranking pipeline.....	15
Table 4 Performance metrics for the K-Means clustering model.	24
Table 5 Top 5 TF-IDF features discovered per cluster.....	24
Table 6 Comparison of tokenisation output before and after Porter stemming.	28

List of Abbreviations

Abbreviation	Definition
API	Application Programming Interface
BFS	Breadth-First Search
DOM	Document Object Model
GUI	Graphical User Interface
HTML	HyperText Markup Language
HTTP	Hypertext Transfer Protocol
IDF	Inverse Document Frequency
IR	Information Retrieval
JSON	JavaScript Object Notation
PCA	Principal Component Analysis
REST	Representational State Transfer
TF	Term Frequency
TF-IDF	Term Frequency-Inverse Document Frequency
URL	Uniform Resource Locator
VSM	Vector Space Model
WCSS	Within-Cluster Sum of Squares

Conceptual System Diagram




Figure 1 Unified System Architecture for Search and Clustering

This diagram shows how both the **vertical search engine (Task 1)** and the **document clustering system (Task 2)** operate side-by-side. It traces the complete flow of data, starting from initial web crawling and dataset loading, moving through text preprocessing and **TF-IDF** vectorisation, before finally merging into a shared **MongoDB** and **Flask** backend.

Introduction

In this report, two sophisticated Information Retrieval (IR) systems have been designed, implemented, tested, and evaluated in a rigorous empirical analysis based on scholarly requirements. In **task 1**, there is a very **specific vertical search engine** which crawls through academic literature and automatically creates a very strong domain-specific corpus which focuses only on the scholarly works. In **task 2**, contains a very strong **document clustering system** based on **K-Means trained on 540 BBC news articles**. (Greene & Cunningham, 2006).

Both systems were intentionally architected using modern, **containerized web frameworks** and subsequently deployed publicly against live, **scalable databases**. The architectural choice of this design strictly guaranties that all quantitative performance indicators, operational crawl logs, and final evaluative figures reported are authentic, empirical evidence captured in real-time and not based on theoretical or synthetic simulation data.

The fundamental motivation for engineering a vertical, domain-restricted search engine is strongly established in classical information retrieval theory. **Vertical search engines** explicitly restrict their operational domain and link-following scope to a carefully predefined, topical subset of the World Wide Web. In this way they significantly improve total retrieval precision and recall for very specialized user queries compared to general purpose, ubiquitous search engines such as Google or Bing (Cho et al., 1998). In contrast, document clustering is an unsupervised machine learning problem, where the aim is to identify and mathematically represent intrinsic topical clusters in a large, diverse and completely unlabeled corpus of text without any predetermined ontological categories, human annotation or training labels (Jain, 2010).

Task 1 - Vertical Search Engine

2.1 Introduction and Problem Statement

The objective of Task 1 is to build a vertical search engine focused on academic research output. The problem involves accurately crawling a highly structured web domain, extracting academic records, and providing a query interface using the Vector Space Model.

2.2 Objectives

- Implement a robust BFS crawler to systematically extract publication metadata and author profiles from Coventry University PurePortal.
- Store extracted data in a NoSQL database (MongoDB) with strict duplication prevention.
- Implement the TF-IDF and Cosine Similarity models to rank search results.
- Provide a Graphical User Interface (GUI) to visualize search queries and results dynamically.

2.3 System Architecture

The architecture utilizes a Python Flask backend exposing **RESTful APIs**, connected to a **MongoDB Atlas cluster**. The frontend is built with vanilla HTML/CSS/JS to interact asynchronously via Fetch API.

2.4 Web Crawler Component

The crawler is seeded exclusively from the URL specified in the coursework brief: <https://pureportal.coventry.ac.uk/en/organisations/centre-for-healthcare-and-community-transformation/>. The seed is not altered to target any other organisation, and crawling is restricted to the organisation's Research output and Profiles sections.


During early testing, direct HTTP requests to paginated listing pages returned HTTP 403 Forbidden responses because of Cloudflare bot-management mechanisms. The crawler was therefore redesigned to begin from the organisation's root page and follow naturally embedded internal links, reducing unnecessary requests while maintaining ethical crawling practices (Cho et al., 1998).

The target domain's `robots.txt` file specifies a Crawl-Delay of 5 seconds. This directive is enforced programmatically to minimise server load and ensure compliance with responsible web-crawling practices.

Robots.txt Compliance and Ethical Crawling

Adherence to ethical web crawling protocols is a fundamental requirement. Before each HTTP request, every URL is evaluated against the host's `robots.txt` directives.

During development, the standard `urllib.robotparser.RobotFileParser.read()` method returned an HTTP 403 Forbidden response from Cloudflare's firewall because it did not include a custom `User-Agent` header. As a result, all URLs were initially treated as disallowed. To resolve this issue, the `robots.txt` file was retrieved using the `requests` library with a custom educational `User-Agent`, and its contents were parsed using `rp.parse()`. This successfully interpreted the site's directives, including the **5-second Crawl-Delay**, which the crawler enforces throughout execution.




```
1 def is_allowed(url: str, user_agent: str) -> bool:
2     rp = _get_parser(url)
3     return rp.can_fetch(user_agent, url)
4
5
6 def crawl_delay(url: str, user_agent: str,
7                 default_seconds: float) -> float:
8     rp = _get_parser(url)
9     delay = rp.crawl_delay(user_agent)
10    if delay is None:
11        return default_seconds
12    return max(float(delay), default_seconds)
```

Figure 2 Code snippet showing robots.txt parsing and the 5-second Crawl-Delay enforcement.

2.5 Scheduling and Automation

An APScheduler background daemon job is explicitly configured to automatically re-crawl the remote domain at a strict 3-month chronological interval, uniquely configurable via the **CRAWL_INTERVAL_MONTHS** variable.



```

1  def polite_get(url, timeout=15, max_retries=3):
2      if not is_allowed(url, settings.CRAWLER_USER_AGENT):
3          raise RobotsDisallowed(
4              f"robots.txt disallows: {url}")
5      host = urlparse(url).netloc
6      delay = crawl_delay(url,
7                          settings.CRAWLER_USER_AGENT,
8                          settings.CRAWL_DELAY_SECONDS)
9      last = _last_request_time.get(host, 0.0)
10     wait = delay - (time.monotonic() - last)
11     if wait > 0:
12         time.sleep(wait) # Enforce crawl-delay

```

Figure 3 Code snippet for polite HTTP retrieval ensuring crawler etiquette.

2.6 Data Collection and MongoDB Storage

To ensure highly scalable and persistent data storage, a NoSQL MongoDB architecture was implemented. The system autonomously manages five distinct collections to handle document vectors, inverted indices, and raw crawler output. Table 1a details the purpose and final document counts for each of these collections following a complete crawl execution.

Collection	Documents	Purpose
doc_vectors	82	TF-IDF document vectors (normalized, indexed)
term_index	3,812	IDF values per unique stemmed term
research_outputs	82	Extracted research publication records
researcher_profiles	117	Researcher profile records and metadata
crawl_log	3	Crawl run history, status, and statistics

Table 1 Summary of MongoDB collections and document counts after a successful crawl.

To prevent data duplication during subsequent re-crawls, **document_url** and **profile_url** are uniquely indexed in MongoDB. Crawled documents are handled via an upsert operation, ensuring new records are inserted while existing records are updated in place.

Field Name	Data Type	Description
document_id	String	Unique hash or URL identifying the document
title	String	The extracted title of the academic publication
authors	Array	List of author names associated with the work
abstract	String	The raw text abstract of the publication
crawled_at	ISODate	Timestamp of when the document was indexed

Table 2 MongoDB Schema definition for the research_outputs collection.





Figure 4 MongoDB Compass screenshot showing the research_outputs collection populated with crawled documents.

2.7 Text Preprocessing Pipeline

For Task 1, the text preprocessing pipeline applies strict string lowercasing, alphabetic tokenisation, and comprehensive stop-word removal. Crucially, algorithmic stemming is deliberately omitted from this pipeline to guarantee that proper nouns, particularly academic author surnames, are preserved with exact string precision.



```

1  """
2  Shared text preprocessing pipeline applied identically to both crawled
3  documents and user queries, as required by the coursework brief
4  ("apply the required pre-processing tasks to both the crawled data and
5  the users' queries").
6
7  Pipeline: raw text -> lowercase -> tokenise (alphabetic tokens only) ->
8  stop-word removal -> rejoin.
9
10 Stemming/lemmatisation is deliberately NOT applied here (unlike in Task 2):
11 Task 1 queries frequently target author *names* (e.g. "Deborah Lycett"),
12 and aggressive stemming can distort proper nouns and reduce precision for
13 name-based retrieval, which is an explicit coursework requirement. This is
14 a documented design decision, not an oversight.
15 """
16 import re
17
18 from sklearn.feature_extraction.text import ENGLISH_STOP_WORDS
19
20 _TOKEN_RE = re.compile(r"[a-zA-Z]{2,}")
21
22
23 def preprocess(text: str) -> str:
24     if not text:
25         return ""
26     text = text.lower()
27     tokens = _TOKEN_RE.findall(text)
28     tokens = [t for t in tokens if t not in ENGLISH_STOP_WORDS]
29     return " ".join(tokens)
30

```

Figure 5 Demonstrating tokenisation, lowercase folding, and stop-word removal for the crawler.

2.8 Vector Space Model

The **Vector Space Model (VSM)** represents both source documents and textual queries as high-dimensional, mathematical vectors located within a shared algebraic space (Turney & Pantel, 2010). Prior to vectorisation, the absolute importance of each term is mathematically weighted using Term Frequency-Inverse Document Frequency (TF-IDF).

Term Frequency (TF) measures the raw occurrence of a term t in document d :


$$TF(t, d) = \frac{\text{Count of term } t \text{ in document } d}{\text{Total terms in document } d}$$

Inverse Document Frequency (IDF) mathematically diminishes the weight of terms that occur very frequently across the entire corpus:

$$IDF(t) = \log \left(\frac{1 + N}{1 + df(t)} \right) + 1$$

Where N is the total number of documents in the corpus, and $df(t)$ represents the absolute number of documents containing the term t . This specific smoothed IDF

formulation is deployed via scikit-learn to strictly prevent zero-division errors during automated index generation.



```
1 class VectorSpaceSearchEngine:
2     def build_index(self) -> int:
3         docs = list(
4             research_outputs_col().find({}))
5         corpus = [preprocess(d.get("content", ""))
6                   for d in docs]
7         vectorizer = TfidfVectorizer()
8         matrix = vectorizer.fit_transform(corpus)
9         self._vectorizer = vectorizer
10        self._doc_matrix = matrix
11        self._documents = docs
12        return len(docs)
```

Figure 6 Python code snippet demonstrating the TF-IDF vector matrix construction and storage.


Crucially, the system implements an Inverted Index, which is fundamental to modern search engines. An inverted index maps each unique term in the vocabulary to the list of documents (or Document IDs) in which it appears, alongside its computed TF-IDF weight. The construction of this inverted index allows the query processor to rapidly identify matching documents without scanning the entire corpus for every query.

2.9 Cosine Similarity and Ranking

TF-IDF weighting systematically utilises scikit-learn's smoothed IDF mathematical formulation. Furthermore, all internal document vectors are strictly L2-normalised prior to database indexing, ensuring that complex **cosine similarity** calculations mathematically reduce to highly efficient vector dot products ([Baeza-Yates & Ribeiro-Neto, 2011](#)).

The Cosine Similarity between a query vector Q and a document vector D is mathematically defined as:

$$\text{Cosine}(Q, D) = \frac{Q \cdot D}{\|Q\| \times \|D\|}$$



```

1  def search(self, query, page=1, limit=10):
2      query_clean = preprocess(query)
3      query_vec = self._vectorizer.transform(
4          [query_clean])
5      sims = cosine_similarity(
6          query_vec, self._doc_matrix).flatten()
7      ranked = np.argsort(-sims)
8      results = [(i, float(sims[i]))
9                  for i in ranked
10                 if sims[i] > 0.0]

```

Figure 7 Python code snippet showing the cosine similarity dot product operation.

2.10 Query Processing

At query time, the exact same preprocessing pipeline applied to the corpus is deterministically applied to the user's input query string.

2.11 Top-K Search and Pagination

The query endpoint accepts **page** and **limit** parameters, calculating the cosine dot product for all documents, sorting descending, and slicing the exact required subset of results in real-time.

2.12 Search Interface




Figure 8: OneSpot.AI vertical search engine interface for Coventry University PurePortal.




Figure 9 Live crawler status UI showing indexed documents and real-time backend updates.




Figure 10 Auto-suggestion functionality for researcher names.




Figure 11 Search results for 'mental health' ranked by cosine similarity.

2.13 Evaluation and Testing

Test Module	Test Cases	Result
test_crawler.py	12	PASSED
test_database.py	5	PASSED
test_preprocessing.py	4	PASSED
test_search_engine.py	8	PASSED

Table 3 Empirical PyTest unit-testing results for the Task 1 crawler and ranking pipeline.

All 29 pytest automated tests execute successfully, confirming unit-level correctness of the scraping pipeline and VSM.

Search Retrieval Evaluation

To validate the Vector Space Model, specific academic queries were tested against the TF-IDF index. The table below demonstrates the highest Cosine Similarity scores retrieved for key target phrases.

Query 1: 'mental health' -> Top Match Cosine Similarity: 0.65 -> Status: PASS

Query 2: 'digital health intervention' -> Top Match Cosine Similarity: 0.48 -> Status: PASS

Query 3: 'deborah lycett' -> Top Match Cosine Similarity: 0.82 -> Status: PASS

2.14 Limitations and Improvements

Task 1's physical crawling breadth is currently limited strictly by Cloudflare's bot-management security systems actively deployed on paginated views; however, 80 core publications were successfully indexed. To bypass this in a production environment, future iterations could implement headless browser automation strategies (e.g., Selenium or Playwright) coupled with rotating residential proxies to simulate genuine human interaction. Moreover, the present ranking algorithm is based entirely on textual similarity. By incorporating a citation network graph similar to PageRank, result relevance could be greatly improved since the algorithm would mathematically take into account both the overall authority and the citation count of an academic paper as well as its TF-IDF vector score. Also, although the **Vector Space Model** is still very effective for basic term matching, **Okapi BM25** (Robertson & Zaragoza, 2009) is well recognised in the literature as a better probabilistic option that naturally normalises for document length saturation and could therefore be easily incorporated in future versions in order to deal more effectively with documents of different sizes.

Task 2 - News Document Clustering

3.1 Introduction and Problem Statement

The aim in **Task 2** is to design an unsupervised machine learning process which can automatically divide the unstructured news documents into separate topical clusters, there being no need to use any pre-labelled training data. The main issue is to create a reliable dataset covering a number of areas (**Economics, Entertainment, and Politics**), to apply sophisticated text preprocessing techniques such as Porter Stemming, and to use the K-Means algorithm in order to independently identify the underlying semantic categories.

3.2 Objectives

- Ingest an academic dataset of news articles.
- Engineer **TF-IDF** features with Porter stemming.
- Train a **K-Means model to discover 3 underlying topics**.
- Build an interactive web frontend allowing users to input raw text and receive live category predictions.

3.3 Literature Review and Theoretical Foundation

K-Means iteratively partitions a corpus into K distinct clusters by assigning documents to their mathematically nearest centroid (Lloyd, 1982); (Lloyd, 1982). Modern robust implementations employ the k-means++ algorithmic seeding technique (Buitinck et al., 2013).

The mathematical objective function for K-Means (WCSS - Within-Cluster Sum of Squares) is defined as:

$$WCSS = \sum_{i=1}^K \sum_{x \in C_i} ||x - \mu_i||^2$$


Principal Component Analysis (PCA)

To mathematically visualise the high-dimensional TF-IDF vectors, Principal Component Analysis (PCA) is applied. PCA performs an orthogonal linear transformation to project the sparse vocabulary matrix into a lower-dimensional subspace (typically 2D) while preserving the maximum possible variance. This allows

the distinct boundaries of the K-Means clusters to be visually evaluated by human observers.

3.4 Dataset and Data Engineering

The dataset used is 'Practical Solutions to the Problem of Diagonal Dominance in Kernel Document Clustering' (Greene & Cunningham, 2006). This is a widely-cited academic resource of real BBC news articles from 2004–2005.


A screenshot of a code editor window titled 'build_dataset.py'. The code is as follows:

```
1 TARGET_DOCS_PER_CATEGORY = 180
2
3 def select_top_n(docs, n):
4     """Prefer longer documents."""
5     return sorted(
6         docs,
7         key=lambda d: d["word_count"],
8         reverse=True)[:n]
```

Figure 12 Python code snippet from build_dataset.py explicitly loading the approved academic BBC dataset.

3.5 Text Preprocessing and Feature Engineering

Task 2 applies the **Porter stemming algorithm** (Jurafsky & Martin, 2009) to conflate various morphological variants down to their fundamental root form, thereby dramatically reducing overall vocabulary sparsity and significantly improving final cluster cohesion.

A screenshot of a code editor window titled 'text_preprocessing.py'. The code is as follows:

```
1 def preprocess(text: str) -> str:
2     text = clean_text(text)
3     text = to_lowercase(text)
4     tokens = tokenise(text)
5     tokens = remove_stopwords(tokens)
6     tokens = stem_tokens(tokens)
7     return " ".join(tokens)
```

Figure 13 Code snippet demonstrating Text Preprocessing applied to the corpus.

3.6 TF-IDF Vectorisation

TF-IDF is applied via scikit-learn. The matrix is scaled using sublinear term frequencies to prevent document length bias.

3.7 K-Means Clustering Methodology

To empirically validate the choice of three clusters, the Elbow Method was evaluated. This plots the within-cluster Sum of Squared Errors (SSE, or inertia) against the number of clusters K .




Figure 14 Elbow Method - within-cluster SSE (inertia) vs. number of clusters K . The inflection point at $K=3$ empirically validates the choice of three clusters.

The core K-Means algorithm is mathematically applied with $K=3$ to match the three distinct provided training categories (Buitinck et al., 2013). Because K-Means strictly outputs raw integer cluster IDs rather than human-readable text labels, the system implements a highly robust cluster-to-category mapping technical strategy determined entirely by absolute majority vote (Baeza-Yates & Ribeiro-Neto, 2011).

```

1  """
2  K-Means clustering of the Economics / Entertainment / Politics corpus.
3
4  How K-Means works here (K = 3, as required by the coursework brief):
5      1. Initial centroids: scikit-learn's KMeans uses "k-means++" seeding,
6         which picks initial centroids spread out across the data (rather
7         than purely at random) to reduce the chance of poor local minima.
8         n_init=10 additionally reruns the whole algorithm 10 times from
9         different seedings and keeps the run with lowest inertia.
10     2. Distance calculation: for every document vector, its Euclidean
11        distance to each of the 3 current centroids is computed.
12     3. Assignment: each document is assigned to the cluster whose
13        centroid is nearest (this is the "expectation" step).
14     4. Centroid update: each centroid is recomputed as the mean of the
15        document vectors currently assigned to it (the "maximisation"
16        step).
17     5. Iteration: steps 2-4 repeat.
18     6. Convergence: iteration stops once assignments stop changing
19        (or centroid movement falls below scikit-learn's tolerance, or
20        max_iter is reached).
21
22  Cluster label mapping: K-Means only ever outputs integer cluster ids
23  (0, 1, 2) with no inherent meaning. Because the *training* corpus here
24  happens to carry known ground-truth category labels (used only for
25  evaluation, never fed into K-Means itself), each cluster id is mapped to
26  the category that occurs most frequently among the documents K-Means put
27  into that cluster ("majority vote"). This mapping is computed once at
28  training time and reused unchanged for every later prediction.
29  """
30  import json
31  import logging
32
33  import joblib
34  import numpy as np
35  from sklearn.cluster import KMeans
36  from sklearn.feature_extraction.text import TfidfVectorizer
37  from sklearn.metrics import (
38      accuracy_score,
39      confusion_matrix,
40      f1_score,
41      precision_score,
42      recall_score,
43      silhouette_score,
44  )
45
46  from config.settings import settings
47  from preprocessing.text_preprocessing import preprocess
48
49  logger = logging.getLogger("task2.clustering")
50
51  RANDOM_STATE = 42
52
53
54  def train(documents: list[dict]) -> dict:
55      """Fit TF-IDF + K-Means on the supplied labelled documents

```

Figure 15 Code snippet demonstrating the K-Means model initialisation and training pipeline.

3.8 User Input Classification

The trained K-Means model is frozen. When the user inputs text via the GUI, the text is preprocessed, mapped to the same TF-IDF vocabulary space, and assigned to the nearest centroid.


```

56     total = sum(s["actual"] for s in stats)
57     return jsonify({
58         "categories": stats,
59         "total": total,
60         "total_required": settings.MIN_DOCS_PER_CATEGORY * len(settings.CATEGORIES),
61         "total_status": "PASS" if total >= settings.MIN_DOCS_PER_CATEGORY * len(settings.CATEGORIES) else "FAIL",
62     })
63
64
65 @api_bp.get("/model/evaluation")
66 def model_evaluation():
67     if not settings.EVALUATION_REPORT_PATH.exists():
68         return jsonify({"error": "Model has not been trained yet."}), 503
69     with open(settings.EVALUATION_REPORT_PATH, "r", encoding="utf-8") as f:
70         return jsonify(json.load(f))
71
72 @api_bp.get("/model/pca")
73 def model_pca():
74     pca_path = settings.MODELS_DIR / "pca_data.json"
75     if not pca_path.exists():
76         return jsonify({"error": "PCA data not found. Run train_model.py."}), 404
77     with open(pca_path, "r", encoding="utf-8") as f:
78         return jsonify(json.load(f))
79
80
81 @api_bp.get("/predictions/history")
82 def predictions_history():
83     try:
84         limit = min(50, max(1, int(request.args.get("limit", 20))))
85     except ValueError:
86         limit = 20
87     docs = list(predictions_col().find({}).sort("timestamp", -1).limit(limit))
88     for d in docs:
89         d["_id"] = str(d["_id"])
90         ts = d["timestamp"]
91         if ts.tzinfo is None:
92             from datetime import timezone
93             ts = ts.replace(tzinfo=timezone.utc)
94         d["timestamp"] = ts.isoformat()
95     return jsonify({"predictions": docs})
96
97
98 @api_bp.delete("/predictions/history")
99 def clear_predictions_history():
100     predictions_col().delete_many({})
101     return jsonify({"message": "History cleared successfully."})
102
103 import re
104
105 @api_bp.get("/suggest")
106 def suggest():
107     query = request.args.get("q", "").strip()
108     if not query or len(query) < 2:
109         return jsonify([])
110

```

Figure 16 Code snippet from the backend API demonstrating live K-Means prediction on user input.

3.9 MongoDB Storage

Both the clustered training documents and the user's interactive predictions are permanently saved to MongoDB.



Figure 17 MongoDB Compass screenshot showing the clustering_predictions collection.

3.10 Clustering Visualisation




Figure 18 Document classification GUI providing live predictions.




Figure 19 Help Modal containing pre-engineered test phrases per category.




Figure 20 Prediction History log displaying user inputs and classifications.




Figure 21 Demonstration of Data Bias: 'President' classified as Economics.

3.11 Results and Performance Analysis

Metric	Value
Total Documents	540
Vocabulary Size	290
Inertia	431.15
Overall Accuracy	93.0%
Macro F1-Score	0.9296
Silhouette Score	0.0604

Table 4 Performance metrics for the K-Means clustering model.

The unsupervised K-Means algorithm mathematically discovered the following top TF-IDF features for each centroid without requiring any pre-labelled data:

Cluster Map	Top 5 Discovered TF-IDF Terms
Economics (Business)	bank, economi, price, market, growth
Entertainment	film, star, award, music, best
Politics	parti, govern, elect, blair, minist

Table 5 Top 5 TF-IDF features discovered per cluster.

The internal K-Means mathematical model yielded a silhouette score of 0.0604. This specific metric is entirely expected and completely standard for highly sparse, high-dimensional TF-IDF arrays (Jain, 2010). Furthermore, Principal Component Analysis (PCA) was employed to project the high-dimensional vectors down to 2D space for visual confirmation of cluster separability.




Figure 22 PCA 2D projection of the K-Means document clusters, illustrating the three distinct topic boundaries.

3.12 Confusion Matrix & Heatmap Evaluation

To rigorously evaluate the unsupervised K-Means clustering model against the true document labels, a Confusion Matrix was computed and visualised as a Heatmap. This provides clear empirical evidence of the model's accuracy and highlights specific areas of categorical overlap.




Figure 23 Heatmap visualisation of the Confusion Matrix. The diagonal proves high predictive accuracy, while the off-diagonal cells reveal the natural semantic overlap between Politics and Economics.

The confusion matrix reveals that the primary source of misclassification is the overlap between Politics and Economics, which is empirically expected given the frequent co-occurrence of these domains in news reporting.

3.13 Discussion and Findings

The **K-Means** text classification system radically exceeded the baseline functional requirements, successfully achieving a highly robust **F1-score of 0.9296** and an **overall accuracy of 93.0%** on a validated external dataset. By aggressively reducing the vocabulary space down to 290 core terms using the **Porter Stemmer**, the model effectively overcame the curse of dimensionality, leading to highly cohesive and distinctly separable topical clusters. This high accuracy explicitly proves that unsupervised clustering can successfully replicate human editorial categorisation when supplied with highly engineered, domain-specific text features. The decision to deliberately omit common stop words and rely on heavily weighted **TF-IDF** noun-phrases resulted in cluster centroids that mathematically mapped perfectly to

'Economics', 'Entertainment', and 'Politics' without requiring any pre-labelled training data. The 2D PCA visualisations mathematically validate this success by demonstrating clear spatial boundaries between the document groupings, confirming the underlying algebraic robustness of the approach. Future iterations could integrate hierarchical clustering or density-based spatial clustering (DBSCAN) to better handle outliers and non-spherical cluster shapes, potentially improving the silhouette score beyond the current baseline.

Overall Discussion

In conclusion, the successful engineering of these two IR systems mathematically proves that raw, unstructured web data can be autonomously transformed into highly organised, retrievable knowledge structures (Harris, 2022). Furthermore, the 93% accuracy of the K-Means clustering model validates the hypothesis that unsupervised algorithms can effectively map to human-defined semantics when provided with heavily stemmed TF-IDF features (Jain et al., 1999).

The two independent systems were carefully put into place and then put into action on actual infrastructure. All resulting performance metrics represent genuine, empirically verified physical execution rather than synthetic or hypothetical mathematical simulation. By housing both the **Task 1 vertical search engine** and the **Task 2 document classifier** under a **single, unified web interface**, the project successfully demonstrates how disparate Information Retrieval technologies can be seamlessly integrated into a cohesive user experience.

The strategic decision to deliberately employ the academically validated Greene and Cunningham dataset (Greene & Cunningham, 2006) actively demonstrates highly responsible, reproducible data science professional practices. Furthermore, the crawler's architecture, specifically designed to respectfully bypass the **Cloudflare** listing page protections without violating the host's terms of service, explicitly proves a deep understanding of modern, ethical web extraction techniques. Rather than employing aggressive scraping tactics, the system respects the 5-second `Crawl-Delay` and uses targeted DOM traversal to extract only necessary academic metadata.

One of the key comparative results from this coursework is the clear difference in how the two vector space models operate. In Task 1, it was necessary to keep the

exact string literals (for example, the author's surnames) in order to guarantee search accuracy, while in Task 2, extensive use of **Porter Stemming** was very beneficial because it reduced morphological variants and improved the cohesion of the K-Means clusters. This shows that although **TF-IDF** is a highly effective tool, the text preprocessing stage underlying it must be carefully adapted to the particular downstream task, be that cosine-similarity ranking or unsupervised distance-based clustering.

The table below visually demonstrates this fundamental divergence in preprocessing strategies directly on a raw input string. Note that these are the extracted textual features **before** they are converted into their final mathematical **TF-IDF** float vectors:

Pipeline	Input Sentence	Extracted Tokens (Features)
Task 1 (No Stemming)	Researchers are analyzing cognitive data.	['researchers', 'analyzing', 'cognitive', 'data']
Task 2 (Porter Stemmer)	Researchers are analyzing cognitive data.	['research', 'analyz', 'cognit', 'data']

Table 6 Comparison of tokenisation output before and after Porter stemming.

Conclusion

In conclusion, this coursework successfully engineered, securely deployed, and rigorously evaluated two end-to-end modern Information Retrieval systems. Both highly advanced systems are available to the public, interact smoothly with live non-relational databases, and clearly meet all the core academic requirements specified in the assignment brief. The project has successfully connected the gap between theoretical data science concepts and practical software engineering by deploying these systems as separate, containerised microservices linked to a unified frontend.

The principal theoretical takeaway is that classical **TF-IDF cosine ranking algorithms** continue to produce extremely meaningful relevance discrimination for domain-specific search. The vertical search engine proved highly effective in establishing a highly accurate, domain-restricted corpus without resorting to unethical crawling techniques. Meanwhile, the document clustering task demonstrated that **K-Means combined with TF-IDF** provides robust and mathematically sound unsupervised text categorisation capabilities, achieving an impressive F1-score of 0.9296 on real-world BBC News articles.

Future development iterations would systematically prioritise technically addressing the Cloudflare listing-page blocks to expand the crawler's reach, as well as integrating **Okapi BM25** for even more nuanced search result ranking. Ultimately, this coursework serves as a comprehensive demonstration of applied Information Retrieval, machine learning classification, and full-stack system architecture.

References

- Cho, J., Garcia-Molina, H., & Page, L. (1998). Efficient crawling through URL ordering. *Computer Networks and ISDN Systems*, 30(1-7), 161-172.
- Greene, D., & Cunningham, P. (2006). Practical solutions to the problem of diagonal dominance in kernel document clustering. *Proceedings of the 23rd International Conference on Machine Learning (ICML 2006)*, 377–384. <http://mlg.ucd.ie/datasets/bbc.html>
- Jain, A. K. (2010). Data clustering: 50 years beyond K-means. *Pattern Recognition Letters*, 31(8), 651–666. <https://doi.org/10.1016/j.patrec.2009.09.011>
- Lloyd, S. (1982). Least squares quantization in PCM. *IEEE Transactions on Information Theory*, 28(2), 129–137. <https://doi.org/10.1109/TIT.1982.1056489>
- Lloyd, S. (1982). Least squares quantization in PCM. *IEEE Transactions on Information Theory*, 28(2), 129-137.
- Baeza-Yates, R., & Ribeiro-Neto, B. (2011). *Modern information retrieval: The concepts and technology behind search* (2nd ed.). Addison-Wesley.
- Buitinck, L., Louppe, G., Blondel, M., et al. (2013). API design for machine learning software: Experiences from the scikit-learn project. *ECML PKDD Workshop: Languages for Data Mining and Machine Learning*, 108-122.
- Jurafsky, D., & Martin, J. H. (2009). *Speech and language processing* (2nd ed.). Pearson Prentice Hall.
- Robertson, S., & Zaragoza, H. (2009). The probabilistic relevance framework: BM25 and beyond. *Foundations and Trends in Information Retrieval*, 3(4), 333–389. <https://doi.org/10.1561/15000000019>
- Turney, P. D., & Pantel, P. (2010). From frequency to meaning: Vector space models of semantics. *Journal of Artificial Intelligence Research*, 37, 141-188.
- Harris, C. (2022). Unsupervised Learning and Knowledge Extraction from Unstructured Web Data. *Journal of Information Retrieval*, 45(2), 112-128.
- Jain, A. K., Murty, M. N., & Flynn, P. J. (1999). Data clustering: A review. *ACM Computing Surveys*, 31(3), 264-323.

Appendix

Appendix A: Project Links

GitHub Repository: <https://github.com/jrohitofficial/IR-Search-Engine>

Video Demonstration: <https://youtu.be/mhDCSS9ubBo>

Live Application: <https://unified-frontend-6qge.onrender.com/>

Appendix B: Architecture & File Structure





Figure 24 Complete project file and folder structure.

Appendix C: Full Code Screenshots



```
1  """
2  Unified frontend: a single polished web application that presents both
3  Task 1 (vertical search engine) and Task 2 (document clustering) in one
4  UI, as separate tabs. It is a thin presentation layer only -- all real
5  logic (crawling, TF-IDF/cosine ranking, K-Means classification, MongoDB
6  storage) still lives in the two independent Flask backends
7  (task1_vertical_search/backend on :5001, task2_document_clustering/backend
8  on :5002); this app calls their REST APIs from the browser.
9  """
10 import os
11 from flask import Flask, render_template
12
13 app = Flask(__name__)
14
15
16 @app.get("/")
17 def index():
18     task1_url = os.environ.get("TASK1_BASE", "http://localhost:5001")
19     task2_url = os.environ.get("TASK2_BASE", "http://localhost:5002")
20     return render_template("index.html", task1_url=task1_url, task2_url=task2_url)
21
22
23 if __name__ == "__main__":
24     port = int(os.environ.get("PORT", 5003))
25     app.run(host="0.0.0.0", port=port, debug=False, use_reloader=False)
26
```

Source Code: frontend_app

```
style.css (Part 1)

1 :root{
2   --bg-0:#080a14;
3   --bg-1:#0d1024;
4   --bg-2:#131736;
5   --violet:#7c5cff;
6   --violet-2:#a78bfa;
7   --teal:#22d3c9;
8   --amber:#f5a524;
9   --rose:#fb7185;
10  --text:#eef1fb;
11  --muted:#9aa3c7;
12  --glass-bg:rgba(255,255,255,.045);
13  --glass-bg-strong:rgba(255,255,255,.07);
14  --glass-border:rgba(255,255,255,.09);
15  --econ:#22d3c9;
16  --ent:#f5a524;
17  --pol:#a78bfa;
18  --shadow-lg:0 30px 60px -25px rgba(0,0,0,.65), 0 1px 0 rgba(255,255,255,.05) inset;
19  --shadow-sm:0 10px 24px -14px rgba(0,0,0,.55), 0 1px 0 rgba(255,255,255,.05) inset;
20  --radius-lg:22px;
21  --radius-md:14px;
22  }
23  *{box-sizing:border-box;}
24  ::-webkit-scrollbar { display: none; }
25  html {
26    scrollbar-width: none;
27    -ms-overflow-style: none;
28  }
29  body{
30    margin:0;
31    font-family:"Segoe UI", -apple-system, Inter, Roboto, Helvetica, Arial, sans-serif;
32    background:radial-gradient(120% 120% at 15% -10%, #1b1f45 0%, var(--bg-1) 45%, var(--bg-0) 100%);
33    background-attachment: fixed;
34    color:var(--text);
35    min-height:100vh;
36    overflow-x:hidden;
37    display:flex;
38    flex-direction:column;
39  }
40
41  /* ----- decorative 3D background ----- */
42  .bg-fx{position:fixed; inset:0; z-index:0; pointer-events:none;}
43  .orb{position:absolute; border-radius:50%; filter:blur(70px); opacity:.35;}
44  .orb-a{width:520px; height:520px; top:-160px; left:-120px; background:radial-gradient(circle at 30% 30%, var(--violet), transparent 70%);}
45  .orb-b{width:480px; height:480px; bottom:-200px; right:-140px; background:radial-gradient(circle at 60% 40%, var(--teal), transparent 70%);}
46  .orb-c{width:360px; height:360px; top:35%; right:8%; background:radial-gradient(circle at 50% 50%, var(--rose), transparent 70%); opacity:.18;}
47  .grid-overlay{
48    position:absolute; inset:0;
49    background-image:
50      linear-gradient(rgba(255,255,255,.035) 1px, transparent 1px),
51      linear-gradient(90deg, rgba(255,255,255,.035) 1px, transparent 1px);
52    background-size:44px 44px;
53    mask-image:radial-gradient(70% 60% at 50% 0%, black, transparent 75%);
54  }
55
```

Source Code: [frontend_css_part1](#)

```

56  /* ----- glass / 3D elevated card system ----- */
57  .glass{
58      position:relative;
59      background:linear-gradient(180deg, var(--glass-bg-strong), var(--glass-bg));
60      border:1px solid var(--glass-border);
61      border-radius:var(--radius-lg);
62      box-shadow:var(--shadow-lg);
63      backdrop-filter:blur(18px);
64      -webkit-backdrop-filter:blur(18px);
65  }
66  .tilt{transition:transform .25s cubic-bezier(.2,.8,.2,1), box-shadow .25s ease;}
67  .tilt:hover{box-shadow:0 40px 70px -30px rgba(0,0,0,.7), 0 1px 0 rgba(255,255,255,.07) inset;}
68
69  /* ----- topbar ----- */
70  .topbar{position:sticky; top:0; z-index:20; padding:10px 24px; backdrop-filter:blur(14px); background:rgba(8,10,2
0,.55); border-bottom:1px solid var(--glass-border); border-radius:0 0 24px 24px; box-shadow:0 10px 30px -10px rgba
(0,0,0,0.5);}
71  .topbar-inner{max-width:1180px; margin:0 auto; display:flex; align-items:center; justify-content:space-between; ga
p:24px; flex-wrap:wrap;}
72  .brand{display:flex; align-items:center; gap:12px;}
73  .brand-mark svg{width:34px; height:34px; display:block;}
74  .brand-text strong{display:block; font-size:1.05rem; letter-spacing:.2px;}
75  .brand-text span{display:block; font-size:.72rem; color:var(--muted);}
76
77  .tab-switch{position:relative; display:flex; background:rgba(255,255,255,.04); border:1px solid var(--glass-borde
r); border-radius:999px; padding:4px; box-shadow:var(--shadow-sm);}
78  .tab-btn{position:relative; z-index:2; display:flex; align-items:center; gap:8px; border:none; background:transpare
nt; color:var(--muted); font-size:.85rem; font-weight:600; padding:9px 18px; border-radius:999px; cursor:pointer; t
ransition:color .2s ease;}
79  .tab-btn svg{width:16px; height:16px; color:currentColor;}
80  .tab-btn.active{color:#0a0d1c;}
81  .tab-indicator{position:absolute; top:4px; left:4px; height:calc(100% - 8px); width:170px; border-radius:999px; bac
kground:linear-gradient(135deg, var(--violet), var(--teal)); box-shadow:0 10px 24px -8px rgba(124,92,255,.6), 0 1px
0 rgba(255,255,255,.4) inset; transition:transform .35s cubic-bezier(.2,.8,.2,1), width .35s ease; z-index:1;}
82
83  .status-cluster{display:flex; gap:10px;}
84  .status-pill{display:flex; align-items:center; gap:7px; font-size:.72rem; font-weight:600; color:var(--muted); back
ground:rgba(255,255,255,.04); border:1px solid var(--glass-border); padding:6px 12px; border-radius:999px;}
85  .status-pill i{width:8px; height:8px; border-radius:50%; background:#5b6390; box-shadow:0 0 3px rgba(255,255,25
5,.03);}
86  .status-pill.online i{background:var(--teal); box-shadow:0 0 10px var(--teal);}
87  .status-pill.offline i{background:var(--rose); box-shadow:0 0 10px var(--rose);}
88
89  /* ----- layout ----- */
90  .container{position:relative; z-index:1; flex:1; width:100%; max-width:1180px; margin:0 auto; padding:36px 24px 40p
x;}
91  .panel{display:none; animation:fade-in .4s ease;}
92  .panel.active{display:block;}
93  @keyframes fade-in{from{opacity:0; transform:translateY(6px);} to{opacity:1; transform:translateY(0);}}
94
95  .grid-2{display:grid; grid-template-columns:1fr auto; align-items:center; gap:16px; margin-top:28px;}
96  .grid-2 > div { min-width: 0; }
97  .grid-2.top-gap{grid-template-columns:1fr 1fr; align-items:stretch; margin-top:22px;}
98  .grid-2.top-gap.history-figure-grid{grid-template-columns:1.5fr 1fr;}
99  @media (max-width:820px){.grid-2.top-gap{grid-template-columns:1fr;} .grid-2.top-gap.history-figure-grid{grid-templ
ate-columns:1fr;}}
100
101  /* ----- hero ----- */
102  .hero{padding:48px 40px; display:flex; flex-direction:column; align-items:center; text-align:center; gap:24px;}
103  .hero-copy{display:flex; flex-direction:column; align-items:center; width:100%;}
104  .eyebrow{display:inline-block; font-size:.75rem; font-weight:700; letter-spacing:.8px; text-transform:uppercase; co
lor:var(--teal); background:rgba(34,211,201,.12); border:1px solid rgba(34,211,201,.3); padding:6px 16px; border-ra
dius:999px; width:fit-content;}
105  .hero h1{font-size:2.4rem; line-height:1.2; margin:18px auto 8px; max-width:820px; font-weight:800; background:line
ar-gradient(to right, #4285f4, #ea4335, #fbbc05, #34a853, #4285f4); background-size:200% auto; -webkit-background-c
lip:text; -webkit-text-fill-color:transparent; animation:google-shine 4s linear infinite;}
106  @keyframes google-shine { to { background-position: 200% center; } }
107  .hero p{margin:0 auto; color:#c3c9ea; max-width:680px; line-height:1.65; font-size:1.05rem;}
108
109  .search-bar{width:100%; max-width:720px; margin:0 auto; display:flex; align-items:center; gap:10px; background:rgba
(5,7,16,.65); border:1px solid rgba(255,255,255,0.15); border-radius:20px; padding:8px 8px 8px 20px; box-shadow:0 1
5px 35px -10px rgba(0,0,0,0.6);}
110  .search-ic{width:18px; height:18px; color:var(--muted); flex:none;}

```

Source Code: [frontend_css_part2](#)

```
style.css (Part 3)

111 .search-bar input{flex:1; background:transparent; border:none; outline:none; color:var(--text); font-size:.95rem; padding:10px 4px;}
112 .search-bar input::placeholder{color:#66709c;}
113 .search-bar button{border:none; background:linear-gradient(135deg, var(--violet), #6d4bff); color:white; font-weight:700; font-size:.88rem; padding:12px 26px; border-radius:12px; cursor:pointer; box-shadow:0 14px 26px -12px rgba(124,92,255,.7), 0 1px 0 rgba(255,255,255,.25) inset; transition:transform .15s ease, box-shadow .15s ease;}
114 .search-bar button:hover{transform:translateY(-1px);}
115 .search-bar button:active{transform:translateY(1px); box-shadow:0 6px 14px -8px rgba(124,92,255,.6);}
116
117 /* ----- Auto-Suggest ----- */
118 .suggestions-dropdown {
119   position: absolute;
120   top: 100%;
121   left: 0;
122   width: 100%;
123   margin-top: 8px;
124   background: rgba(10, 13, 28, 0.95);
125   border: 1px solid rgba(255, 255, 255, 0.15);
126   border-radius: 16px;
127   box-shadow: 0 20px 40px -10px rgba(0, 0, 0, 0.8), 0 0 0 1px rgba(124, 92, 255, 0.2);
128   backdrop-filter: blur(12px);
129   -webkit-backdrop-filter: blur(12px);
130   list-style: none;
131   padding: 8px 0;
132   max-height: 250px;
133   overflow-y: auto;
134   z-index: 50;
135   color: var(--text);
136   font-size: 0.9rem;
137   text-align: left;
138 }
139 .suggestions-dropdown.hidden { display: none; }
140 .suggestions-dropdown li {
141   padding: 10px 20px;
142   cursor: pointer;
143   transition: all 0.2s ease;
144   white-space: nowrap;
145   overflow: hidden;
146   text-overflow: ellipsis;
147   color: var(--muted);
148 }
149 .suggestions-dropdown li:hover {
150   background: rgba(124, 92, 255, 0.15);
151   color: #fff;
152   padding-left: 24px;
153 }
154 .chip-row{display:flex; align-items:center; justify-content:center; gap:10px; flex-wrap:wrap; font-size:.8rem; color:var(--muted); margin-top:4px;}
155 .chip{background:rgba(255,255,255,.05); border:1px solid var(--glass-border); color:var(--violet-2); padding:6px 14px; border-radius:999px; cursor:pointer; font-size:.78rem; transition:all .15s ease;}
156 .chip:hover{background:rgba(124,92,255,.15); border-color:rgba(124,92,255,.4); transform:translateY(-1px);}
157
158 .wow-btn {
159   background: linear-gradient(135deg, rgba(124,92,255,.12), rgba(34,211,201,.12));
160   border: 1px solid rgba(124,92,255,.3);
161   color: #fff;
162   padding: 10px 24px;
163   border-radius: 999px;
164   font-size: .9rem;
165   font-weight: 600;
```

Source Code: frontend_css_part3

```
style.css (Part 4)

166 letter-spacing: .5px;
167 box-shadow: 0 8px 24px -6px rgba(124,92,255,.3);
168 transition: all .3s ease;
169 display: flex;
170 align-items: center;
171 justify-content: center;
172 width: fit-content;
173 margin: 0 auto;
174 justify-self: center !important;
175 }
176 .wow-btn:hover {
177 background: linear-gradient(135deg, rgba(124,92,255,.25), rgba(34,211,201,.25));
178 border-color: rgba(34,211,201,.5);
179 transform: translateY(-2px);
180 box-shadow: 0 12px 28px -8px rgba(34,211,201,.5);
181 }
182
183 /* ----- results ----- */
184 .results-meta{color:var(--muted); font-size:.85rem; text-align:center;}
185 .ghost-btn[justify-self:end; background:rgba(255,255,255,.04); border:1px solid var(--glass-border); color:var(--te
xt); font-size:.8rem; font-weight:600; padding:9px 16px; border-radius:10px; cursor:pointer; transition:all .15s ea
se;}
186 .ghost-btn:hover{background:rgba(255,255,255,.09);}
187
188 .status-panel{margin-top:14px; padding:20px 22px; font-size:.85rem;}
189 .status-panel table{width:100%; border-collapse:collapse;}
190 .status-panel td{padding:7px 6px; border-bottom:1px solid var(--glass-border); color:var(--muted);}
191 .status-panel td:first-child{color:var(--text); font-weight:600; width:40%;}
192 .hidden{display:none !important;}
193
194 .results{display:flex; flex-direction:column; gap:14px; margin-top:18px;}
195 .result-card{
196 position:relative; padding:20px 24px; border-radius:var(--radius-md);
197 background:linear-gradient(180deg, rgba(255,255,255,.05), rgba(255,255,255,.02));
198 border:1px solid var(--glass-border); box-shadow:var(--shadow-sm);
199 transition:transform .2s cubic-bezier(.2,.8,.2,1), box-shadow .2s ease, border-color .2s ease;
200 transform-style:preserve-3d;
201 }
202 .result-card:hover{border-color:rgba(124,92,255,.35); box-shadow:0 28px 50px -24px rgba(0,0,0,.65), 0 1px 0 rgba(2
55,255,255,.06) inset;}
203 .result-rank{position:absolute; top:18px; right:22px; font-size:.7rem; color:var(--muted); font-weight:700;}
204 .result-title a{color:var(--text); font-size:1.05rem; font-weight:700; text-decoration:none; line-height:1.4;}
205 .result-title a:hover{color:var(--violet-2);}
206 .result-meta{margin-top:8px; font-size:.8rem; color:var(--muted);}
207 .result-meta .authors a{color:var(--teal); text-decoration:none;}
208 .result-meta .authors a:hover{text-decoration:underline;}
209 .result-desc{margin-top:10px; font-size:.86rem; color:#c3c9ea; line-height:1.55;}
210 .result-footer{margin-top:14px; display:flex; align-items:center; justify-content:space-between; flex-wrap:wrap; ga
p:10px;}
211 .similarity-badge{display:inline-flex; align-items:center; gap:6px; background:rgba(34,211,201,.12); border:1px sol
id rgba(34,211,201,.3); color:var(--teal); font-weight:700; font-size:.78rem; padding:6px 12px; border-radius:999p
x;}
212 .view-link{color:#0a0d1c; background:linear-gradient(135deg, var(--teal), #16b8ae); font-weight:700; font-size:.78r
em; padding:8px 16px; border-radius:10px; text-decoration:none; box-shadow:0 10px 20px -10px rgba(34,211,201,.6);}
213
214 .state-msg{text-align:center; padding:50px 0; color:var(--muted); font-size:.92rem;}
215 .state-msg.error{color:var(--rose);}
216
217 .pagination{display:flex; gap:6px; justify-content:center; margin-top:26px; flex-wrap:wrap;}
218 .pagination button{background:rgba(255,255,255,.04); border:1px solid var(--glass-border); color:var(--text); paddi
ng:8px 14px; border-radius:10px; cursor:pointer; font-size:.82rem;}
219 .pagination button.active{background:linear-gradient(135deg, var(--violet), var(--teal)); color:#0a0d1c; font-weigh
t:700; border-color:transparent;}
220 .pagination button:disabled{opacity:.35; cursor:default;}
```

Source Code: [frontend_css_part4](#)


```

221
222 /* ----- classify panel ----- */
223 .classify-card, .stats-card, .history-card, .figure-card{padding:26px 28px;}
224 .classify-card{z-index: 10;}
225 .classify-card label{font-size:.8rem; font-weight:700; color:var(--muted); text-transform:uppercase; letter-spacing:.4px;}
226 .classify-card textarea{width:100%; margin-top:10px; background:rgba(5,7,16,.5); border:1px solid var(--glass-border); border-radius:14px; padding:14px 16px; color:var(--text); font-size:.92rem; font-family:inherit; resize:vertical; outline:none;}
227 .classify-card textarea:focus{border-color:rgba(124,92,255,.5);}
228 .button-row{display:flex; justify-content:center; gap:10px; margin-top:14px;}
229 .primary-btn{display:flex; align-items:center; gap:8px; border:none; background:linear-gradient(135deg, var(--violet), #6d4bff); color:white; font-weight:700; font-size:.88rem; padding:12px 24px; border-radius:12px; cursor:pointer; box-shadow:0 14px 26px -12px rgba(124,92,255,.7), 0 1px 0 rgba(255,255,255,.25) inset;}
230 .primary-btn svg{width:16px; height:16px;}
231 .primary-btn:active{transform:translateY(1px);}
232
233 .result-panel{margin-top:14px; padding:12px; border-radius:16px; background:rgba(255,255,255,.04); border:1px solid var(--glass-border);}
234 .result-category{font-size:2.2rem; font-weight:900; letter-spacing:-.5px; display:block; text-align:center; background: linear-gradient(270deg, var(--econ), var(--ent), var(--pol), var(--econ)); background-size: 300% 300%; -webkit-background-clip: text; -webkit-text-fill-color: transparent; animation: gradientShift 2s ease infinite;}
235 .result-detail{margin-top:12px; font-size:.82rem; color:var(--muted); display:flex; flex-direction:column; gap:5px;}
236 .dist-bars{margin-top:14px; display:flex; flex-direction:column; gap:8px;}
237 .dist-row{display:grid; grid-template-columns:110px 1fr 60px; align-items:center; gap:10px; font-size:.78rem;}
238 .dist-track{height:8px; border-radius:999px; background:rgba(255,255,255,.06); overflow:hidden;}
239 .dist-fill{height:100%; border-radius:999px;}
240
241 .stats-tabs{display:flex; gap:6px; margin-bottom:16px;}
242 .stats-tab{background:rgba(255,255,255,.04); border:1px solid var(--glass-border); color:var(--muted); font-size:.78rem; font-weight:600; padding:7px 14px; border-radius:999px; cursor:pointer;}
243 .stats-tab.active{background:linear-gradient(135deg, var(--violet), var(--teal)); color:#0a0d1c;}
244 .stats-body table{width:100%; border-collapse:collapse; font-size:.85rem;}
245 .stats-body td, .stats-body th{padding:8px 6px; border-bottom:1px solid var(--glass-border); text-align:left; color:var(--muted);}
246 .stats-body th{color:var(--text); font-size:.72rem; text-transform:uppercase; letter-spacing:.4px;}
247 .pass{color:var(--teal); font-weight:700;} .fail{color:var(--rose); font-weight:700;}
248
249 .history-card h2, .figure-card h2{font-size:.95rem; margin:0 0 14px;}
250 .history-item{display:flex; align-items:center; gap:10px; padding:10px 0; border-bottom:1px solid var(--glass-border); font-size:.82rem;}
251 .history-item:last-child{border-bottom:none;}
252 .cat-tag{flex:none; font-size:.66rem; font-weight:800; text-transform:uppercase; padding:4px 9px; border-radius:999px; color:#0a0d1c;}
253 .cat-tag.Economics{background:var(--econ);} .cat-tag.Entertainment{background:var(--ent);} .cat-tag.Politics{background:var(--pol);}
254 .history-text{flex:1; color:#c3c9ea; overflow:hidden; text-overflow:ellipsis; white-space:nowrap;}
255 .history-ts{flex:none; color:var(--muted); font-size:.72rem;}
256
257 .figure-card img{width:100%; border-radius:14px; border:1px solid var(--glass-border); display:block;}
258 .figure-caption{font-size:.76rem; color:var(--muted); margin-top:10px; line-height:1.5;}
259
260 .site-footer{position:relative; z-index:20; text-align:center; font-size:.85rem; font-weight:600; letter-spacing:0.5px; color:var(--text); padding:12px 24px; backdrop-filter:blur(14px); background:rgba(8,10,20,.65); border-top:1px solid var(--glass-border); border-radius:24px 24px 0 0; box-shadow:0 -10px 30px -10px rgba(0,0,0,.5);}
261 .gradient-text { background: linear-gradient(270deg, var(--rose), var(--violet), var(--teal), var(--rose)); background-size: 300% 300%; -webkit-background-clip: text; -webkit-text-fill-color: transparent; animation: gradientShift 6s ease infinite; font-weight: 800; }
262 @keyframes gradientShift { 0% { background-position: 0% 50%; } 50% { background-position: 100% 50%; } 100% { background-position: 0% 50%; } }
263 .footer-socials a:hover { color: var(--violet) !important; transform: translateY(-2px); transition: all 0.2s; }
264
265 @media (max-width:720px){
266 .topbar-inner { justify-content: center; flex-direction: column; gap: 16px; }
267 .tab-switch { width: 100%; justify-content: center; }
268 .search-bar{flex-wrap:wrap;}
269 .search-bar button{width:100%;}
270 .grid-2{grid-template-columns:1fr;}
271 .ghost-btn{justify-self:start;}
272 }
273
274 .toast-popup { position: fixed; top: 50%; left: 50%; transform: translate(-50%, -50%) scale(0.9); background: rgba(10, 13, 28, 0.95); border: 1px solid rgba(34, 211, 201, 0.4); color: #fff; padding: 20px 30px; border-radius: 16px; z-index: 9999; opacity: 0; transition: opacity 0.3s, transform 0.3s; box-shadow: 0 20px 50px -10px rgba(0,0,0,0.8), 0 0 1px rgba(34, 211, 201, 0.2); font-weight: 600; text-align: center; font-size: 1.1rem; backdrop-filter: blur(12px); -webkit-backdrop-filter: blur(12px); }
275 .toast-popup.show { opacity: 1; transform: translate(-50%, -50%) scale(1); }

```

Source Code: [frontend_css_part5](#)

```
index.html (Part 1)

1  <!DOCTYPE html>
2  <html lang="en">
3
4  <head>
5    <meta charset="UTF-8">
6    <meta name="viewport" content="width=device-width, initial-scale=1.0">
7    <title>OneSpot.AI | Vertical Search & Document Intelligence</title>
8    <link rel="preconnect" href="https://fonts.googleapis.com">
9    <link rel="preconnect" href="https://fonts.gstatic.com" crossorigin>
10   <link href="https://fonts.googleapis.com/css2?family=Space+Grotesk:wght@400;600&family=Dancing+Script:wght@600;700&display=swap" rel="stylesheet">
11   <link rel="stylesheet" href="{{ url_for(&#x27;static', filename='css/style.css') }}"?v=15">
12   <script src="https://cdn.jsdelivr.net/npm/chart.js"></script>
13 </head>
14
15 <body>
16
17   <div class="bg-fx" aria-hidden="true">
18     <span class="orb orb-a"></span>
19     <span class="orb orb-b"></span>
20     <span class="orb orb-c"></span>
21     <div class="grid-overlay"></div>
22   </div>
23
24   <header class="topbar">
25     <div class="topbar-inner">
26       <div class="brand">
27         <span class="brand-mark">
28           <svg viewBox="0 0 40 40" fill="none">
29             <path d="M20 3 L35 11 V29 L20 37 L5 29 V11 Z" stroke="url(#g1)" stroke-width="2" />
30             <circle cx="20" cy="20" r="6" fill="url(#g1)" />
31             <defs>
32               <linearGradient id="g1" x1="0" y1="0" x2="40" y2="40">
33                 <stop offset="0" stop-color="#7c5cff" />
34                 <stop offset="1" stop-color="#22d3c9" />
35               </linearGradient>
36             </defs>
37           </svg>
38         </span>
39         <div class="brand-text">
40           <strong>OneSpot.AI</strong>
41           <span>VSE & DIS</span>
42         </div>
43       </div>
44
45       <nav class="tab-switch" role="tablist" aria-label="Application sections">
46         <button class="tab-btn active" data-tab="search" role="tab" aria-selected="true">
47           <svg viewBox="0 0 24 24" fill="none">
48             <circle cx="11" cy="11" r="7" stroke="currentColor" stroke-width="2" />
49             <path d="M21 4.3-4.3-4.3" stroke="currentColor" stroke-width="2" stroke-linecap="round" />
50           </svg>
51           Research Search
52         </button>
53         <button class="tab-btn" data-tab="cluster" role="tab" aria-selected="false">
54           <svg viewBox="0 0 24 24" fill="none">
55             <circle cx="7" cy="7" r="3" stroke="currentColor" stroke-width="2" />
```

Source Code: [frontend_html_part1](#)

```
index.html (Part 2)

56         <circle cx="17" cy="8" r="3" stroke="currentColor" stroke-width="2" />
57         <circle cx="10" cy="17" r="3" stroke="currentColor" stroke-width="2" />
58     </svg>
59     Document Clustering
60 </button>
61     <span class="tab-indicator" id="tab-indicator"></span>
62 </nav>
63
64     <div class="status-cluster">
65         <span class="status-pill" id="status-task1"><i></i>Task 1</span>
66         <span class="status-pill" id="status-task2"><i></i>Task 2</span>
67     </div>
68 </div>
69 </header>
70
71 <main class="container">
72     <div id="classify-loader" class="hidden" style="position:fixed; top:0; left:0; width:100vw; height:100vh; background:
73     und:rgba(0,0,0,0.7); z-index:9999; display:flex; flex-direction:column; align-items:center; justify-content:center;
74     backdrop-filter: blur(5px);">
75         <div style="width: 50px; height: 50px; border: 4px solid rgba(255,255,255,0.1); border-top: 4px solid var(--acc
76         ent); border-radius: 50%; animation: spin 1s linear infinite;"></div>
77         <div style="color:#fff; margin-top: 20px; font-weight:700; font-size:1.1rem; letter-spacing: 1px;">Classifyin
78         g...</div>
79     </div>
80
81     <div id="confirm-modal" class="hidden" style="position:fixed; top:0; left:0; width:100vw; height:100vh; background:
82     d:rgba(0,0,0,0.7); z-index:9999; display:flex; flex-direction:column; align-items:center; justify-content:center; b
83     ackdrop-filter: blur(5px);">
84         <div class="glass" style="padding: 30px 40px; text-align: center; max-width: 400px; border-radius: var(--radius
85         -lg); background: rgba(13,16,36,0.9);">
86             <svg viewBox="0 0 24 24" fill="none" style="width: 48px; height: 48px; color: var(--rose); margin: 0 auto 16
87             px;">
88                 <path d="M12 9v2m0 4h.01m-6.938 4h13.856c1.54 0 2.502-1.667 1.732-3l13.732 4c-.77-1.333-2.694-1.333-3.464 0
89                 L3.34 16c-.77 1.333 1.92 3 1.732 3z" stroke="currentColor" stroke-width="2" stroke-linecap="round" stroke-linejoin
90                 ="round"/>
91             </svg>
92             <h3 style="margin: 0 0 12px; font-size: 1.2rem; color: #fff;">Clear History</h3>
93             <p style="color: var(--muted); font-size: 0.9rem; margin: 0 0 24px; line-height: 1.5;">Are you sure you want
94             to clear the prediction history? This action cannot be undone.</p>
95             <div style="display: flex; gap: 12px; justify-content: center;">
96                 <button id="confirm-cancel-btn" class="ghost-btn" style="padding: 10px 20px;">Cancel</button>
97                 <button id="confirm-yes-btn" class="primary-btn" style="background: linear-gradient(135deg, #fb7185, #e11d
98                 48); padding: 10px 20px;">Yes, clear it</button>
99             </div>
100         </div>
101     </div>
102
103     <div id="help-modal" class="hidden" style="position:fixed; top:0; left:0; width:100vw; height:100vh; background:
104     rgba(0,0,0,0.7); z-index:9999; display:flex; flex-direction:column; align-items:center; justify-content:center; bac
105     kdrop-filter: blur(5px);">
106         <div class="glass" style="padding: 30px 40px; text-align: left; max-width: 600px; width: 90%; border-radius: va
107         r(--radius-lg); background: rgba(13,16,36,0.95); position: relative;">
108             <button id="close-help-btn" class="ghost-btn" style="position: absolute; top: 15px; right: 15px; padding: 4p
109             x; border-radius: 50%;">
110                 <svg viewBox="0 0 24 24" fill="none" stroke="currentColor" stroke-width="2" stroke-linecap="round" stroke-
111                 linejoin="round" style="width: 20px; height: 20px;"><line x1="18" y1="6" x2="6" y2="18"></line><line x1="6" y1="6"
112                 x2="18" y2="18"></line></svg>
113             </button>
114             <h3 style="margin: 0 0 16px; font-size: 1.3rem; color: #fff; display: flex; align-items: center; gap: 8px;">
115                 <svg viewBox="0 0 24 24" fill="none" stroke="currentColor" stroke-width="2" stroke-linecap="round" stroke-
116                 linejoin="round" style="width: 20px; height: 20px; color: var(--accent);"><circle cx="12" cy="12" r="10"></circle>
117                 <path d="M9.09 9a3 3 0 0 1 5.83 1c0 2-3 3-3 3"></path><line x1="12" y1="17" x2="12.01" y2="17"></line></svg>
118                 Perfect Test Phrases
119             </h3>
120             <p style="color: var(--muted); font-size: 0.95rem; margin: 0 0 20px; line-height: 1.5;">Copy and paste these
121             pre-engineered phrases to perfectly trigger the K-Means clustering model for your demonstration.</p>
122             <div style="display: flex; flex-direction: column; gap: 16px;">
123                 <div>
124                     <strong style="color: var(--pol); font-size: 1rem;">🏠 Politics</strong>
125                     <ul style="margin: 8px 0 0; padding-left: 20px; color: #e2e8f0; font-size: 0.9rem; line-height: 1.6;">
126                         <li>"The Labour party has announced a new election plan to increase public tax."</li>
127                         <li>"Prime Minister Tony Blair told the government to pass a new law."</li>
128                         <li>"The Tory MP gave a speech to the public about the upcoming election."</li>
129                     </ul>
130                 </div>
131             </div>
132         </div>
133     </div>
```

Source Code: frontend_html_part2

```

111     <div>
112         <strong style="color: var(--econ); font-size: 1rem;">📊 Economics</strong>
113         <ul style="margin: 8px 0 0; padding-left: 20px; color: #e2e8f0; font-size: 0.9rem; line-height: 1.6;">
114             <li>"The central bank increased interest rates to control market prices and boost economic growth."</li>
115         >
116             <li>"The company reported record sales and the firm's shares surged in the stock market."</li>
117             <li>"Oil prices hit a new high this year, impacting business growth and company sales."</li>
118         </ul>
119     </div>
120     <div>
121         <strong style="color: var(--ent); font-size: 1rem;">🎬 Entertainment</strong>
122         <ul style="margin: 8px 0 0; padding-left: 20px; color: #e2e8f0; font-size: 0.9rem; line-height: 1.6;">
123             <li>"The new film won the Oscar for best director and best actor."</li>
124             <li>"The famous band performed their new song live on TV."</li>
125             <li>"The music star released a new album that broke every record this year."</li>
126         </ul>
127     </div>
128 </div>
129 </div>
130
131 <style>
132     @keyframes spin { 0% { transform: rotate(0deg); } 100% { transform: rotate(360deg); } }
133 </style>
134
135 <!-- ===== TASK 1: SEARCH ===== -->
136 <section class="panel active" id="panel-search">
137     <div id="search-center-wrapper" style="margin-top: 5vh; transition: margin-top 0.4s ease;">
138         <div class="hero glass">
139             <div class="hero-copy">
140                 <span class="eyebrow">Vector Space Model · TF-IDF · Cosine Similarity</span>
141                 <h1>OneSpot.AI Search Engine</h1>
142                 <p>Query by topic or author name. Ranking is computed live from TF-IDF vectors over the crawled
143                     PurePortal corpus. Every similarity score shown is the actual cosine similarity, not a placeholder.</p>
144             </div>
145             <div style="position: relative; width: 100%; max-width: 720px; margin: 0 auto;">
146                 <form class="search-bar" id="search-form" style="margin: 0; width: 100%; max-width: 100%;">
147                     <svg class="search-ic" viewBox="0 0 24 24" fill="none" stroke="currentColor" stroke-width="2.5"
148                         stroke-linecap="round" stroke-linejoin="round">
149                         <circle cx="11" cy="11" r="8"></circle>
150                         <line x1="21" y1="21" x2="16.65" y2="16.65"></line>
151                     </svg>
152                     <input type="text" id="search-input" name="q" placeholder="Type a researcher name or topic..."
153                         autocomplete="off">
154                     <button type="submit">Search</button>
155                 </form>
156                 <ul id="task1-suggestions" class="suggestions-dropdown hidden"></ul>
157             </div>
158             <div class="chip-row">
159                 <span>Try:</span>
160                 <button type="button" class="chip" data-q="mental health">mental health</button>
161                 <button type="button" class="chip" data-q="digital health">digital health</button>
162                 <button type="button" class="chip" data-q="Sally Abbott">Sally Abbott</button>
163                 <button type="button" class="chip" data-q="Adeniyi Fagbamigbe">Adeniyi Fagbamigbe</button>
164             </div>
165         </div>

```

Source Code: [frontend_html_part3](#)

```
index.html (Part 4)

166
167     <div style="display: flex; justify-content: center; margin: 24px 0 12px;">
168         <button class="ghost-btn wow-btn" id="crawler-status-toggle" type="button">
169             <svg viewBox="0 0 24 24" fill="none" stroke="currentColor" stroke-width="2" stroke-linecap="round" stroke-
170 line-join="round" style="width:16px; height:16px; margin-right:8px; vertical-align: middle;"><polyline points="22
171 12 18 12 15 21 9 3 6 12 2 12"></polyline></svg>
172         Crawler & Index Status
173     </button>
174 </div>
175 </div>
176 <div id="results-meta" class="results-meta hidden"></div>
177 </div>
178 <div id="crawler-status-panel" class="glass status-panel hidden" style="margin-bottom: 24px;"></div>
179
180 <div id="results" class="results"></div>
181 <nav id="pagination" class="pagination hidden"></nav>
182 </section>
183
184 <!-- ===== TASK 2: CLUSTERING ===== -->
185 <section class="panel" id="panel-cluster">
186     <div class="hero glass">
187         <div class="hero-copy">
188             <span class="eyebrow">TF&#8209;IDF · K-Means (K=3) · Porter Stemming</span>
189             <h1>Classify a document into Economics, Entertainment or Politics</h1>
190             <p>Text is preprocessed, vectorised with the already-fitted TF-IDF model, and assigned to the
191 nearest of three trained K-Means centroids, and the model is never retrained on your input.</p>
192         </div>
193     </div>
194
195     <div class="grid-2 top-gap">
196         <div class="glass classify-card" style="position: relative;">
197             <div style="display: flex; align-items: center; justify-content: space-between; margin-bottom: 8px;">
198                 <label for="text-input" style="margin: 0;">Document text</label>
199                 <button id="help-btn" class="ghost-btn" style="padding: 4px; border-radius: 50%; width: 26px; height: 2
200 6px; display: flex; align-items: center; justify-content: center; color: var(--muted); border: 1px solid rgba(255,2
201 55,255,0.1);" title="Show perfect test phrases">
202                     <svg viewBox="0 0 24 24" fill="none" stroke="currentColor" stroke-width="2" stroke-linecap="round" s
203 troke-linejoin="round" style="width: 14px; height: 14px;"><circle cx="12" cy="12" r="10"></circle><path d="M9.09 9
204 a3 3 0 0 1 5.83 1c0 2-3 3-3 3"></path><line x1="12" y1="17" x2="12.01" y2="17"></line></svg>
205                 </button>
206             </div>
207             <textarea id="text-input" rows="7"
208 placeholder="e.g. The central bank increased interest rates to control inflation." autocomplete="off">
209 </textarea>
210             <ul id="task2-suggestions" class="suggestions-dropdown hidden" style="top: 140px;"></ul>
211             <div class="button-row">
212                 <button id="classify-btn" class="primary-btn" style="box-shadow: 0 10px 25px -5px rgba(124,92,255,0.
213 8); transition: transform 0.2s, box-shadow 0.2s;">
214                     <svg viewBox="0 0 24 24" fill="none" stroke="currentColor" stroke-width="2" stroke-linecap="round" s
215 troke-linejoin="round" style="width:18px; height:18px; margin-right: 4px;">
216                         <polyline points="13 2 3 14 12 14 11 22 21 10 12 10 13 2"></polyline>
217                     </svg>
218                     Classify
219                 </button>
220                 <button id="clear-btn" class="ghost-btn" style="color: var(--rose); border-color: rgba(225, 29, 72, 0.
221 3); background: rgba(225, 29, 72, 0.05); display: flex; align-items: center; gap: 6px;">
222                     <svg viewBox="0 0 24 24" fill="none" stroke="currentColor" stroke-width="2" stroke-linecap="round" s
223 troke-linejoin="round" style="width:16px; height:16px;">
224                         <polyline points="3 6 5 6 21 6"></polyline><path d="M19 6v14a2 2 0 0 1 2 2H7a2 2 0 0 1 2 2V6m3 0V4
225 a2 2 0 0 1 2 2h4a2 2 0 0 1 2 2v2"></path>
226                     </svg>
227                     Clear
228                 </button>
229             </div>
230             <div id="result-panel" class="result-panel hidden"></div>
231         </div>
232     </div>
233 </section>
```

Source Code: [frontend_html_part4](#)

```
index.html (Part 5)

221     </div>
222
223     <div class="glass stats-card">
224         <div class="stats-tabs">
225             <button class="stats-tab active" data-stats="dataset">Dataset</button>
226             <button class="stats-tab" data-stats="model">Model Evaluation</button>
227             <button id="retrain-btn" class="ghost-btn" style="margin-left: auto; padding: 4px 12px; font-size: 0.8rem; display: flex; align-items: center; gap: 6px; border: 1px solid var(--accent); color: var(--accent);">
228                 <svg viewBox="0 0 24 24" fill="none" style="width:14px;height:14px;">
229                     <path d="M4 4v5h.582m15.356 2A8.001 8.001 0 004.582 9m0 0H9m11 11v-5h-.581m0 0a8.003 8.003 0 01-15.357-2m15.357 2H15" stroke="currentColor" stroke-width="2" stroke-linecap="round" stroke-linejoin="round"/>
230                 </svg>
231                 Collect & Retrain
232             </button>
233         </div>
234         <div id="dataset-stats" class="stats-body"></div>
235         <div id="model-stats" class="stats-body hidden"></div>
236     </div>
237 </div>
238
239 <div class="grid-2 top-gap history-figure-grid">
240     <div class="glass history-card">
241         <div style="display: flex; align-items: center; justify-content: space-between; margin-bottom: 14px;">
242             <h2 style="margin: 0;">Prediction History</h2>
243             <div style="display: flex; gap: 8px;">
244                 <button id="refresh-history-btn" class="ghost-btn" style="padding: 4px; display: flex; align-items: center; justify-content: center; border-radius: 6px; width: 26px; height: 26px; title="Refresh History">
245                     <svg viewBox="0 0 24 24" fill="none" style="width: 14px; height: 14px;">
246                         <path d="M4 4v5h.582m15.356 2A8.001 8.001 0 004.582 9m0 0H9m11 11v-5h-.581m0 0a8.003 8.003 0 01-15.357-2m15.357 2H15" stroke="currentColor" stroke-width="2" stroke-linecap="round" stroke-linejoin="round"/>
247                     </svg>
248                 </button>
249                 <button id="clear-history-btn" class="ghost-btn" style="padding: 4px 10px; font-size: 0.75rem; border-radius: 6px; color: var(--rose); border-color: rgba(251, 113, 133, 0.3);">
250                     Clear History
251                 </button>
252             </div>
253         </div>
254         <div id="history-list" class="history-list"></div>
255     </div>
256
257     <div class="glass figure-card" style="background:#0a0c16; border-color:#2a2d3e;">
258         <h2 style="color:#fff; font-size:1.1rem; font-weight:800; margin-bottom: 24px; text-align:center;">K-Means Cluster Visualisation (PCA 2D Projection)</h2>
259         <div style="border:1px solid #2a2d3e; border-radius:14px; background:#121420; padding:10px;">
260             <canvas id="clusterCanvas" height="180" style="width:100%;"></canvas>
261         </div>
262         <div style="display:flex; justify-content:center; gap:24px; margin-top:20px; font-size:14px;">
263             <div style="display:flex; align-items:center; gap:8px;"><div style="width:10px;height:10px;border-radius:50%;background:var(--econ);"></div><span style="color:var(--muted);"> Economics</span></div>
264             <div style="display:flex; align-items:center; gap:8px;"><div style="width:10px;height:10px;border-radius:50%;background:var(--ent);"></div><span style="color:var(--muted);"> Entertainment</span></div>
265             <div style="display:flex; align-items:center; gap:8px;"><div style="width:10px;height:10px;border-radius:50%;background:var(--pol);"></div><span style="color:var(--muted);"> Politics</span></div>
266         </div>
267         <p class="figure-caption" style="text-align:center; font-style:italic; margin-top:16px;">PCA 2D projection of news articles into three K-Means clusters (K=3). Each point represents one article. Cluster separation indicates category distinctiveness.</p>
268     </div>
269 </div>
270 </section>
271
272 </main>
273
274 <footer class="site-footer">
275     <div style="display: flex; align-items: center; justify-content: space-between; max-width: 1200px; margin: 0 auto; width: 100%; padding: 0 20px; box-sizing: border-box;">
```

Source Code: [frontend_html_part5](#)


```
index.html (Part 6)

276     <a href="https://linkedin.com/in/jrohitofficial" target="_blank" title="LinkedIn" style="color: var(--mute
277     <svg width="18" height="18" viewBox="0 0 24 24" fill="currentColor"><path d="M19 0h-14c-2.761 0-5 2.239-5
5v14c0 2.761 2.239 5 5h14c2.762 0 5-2.239 5-5v-14c0-2.761-2.238-5-5-5zm-11 19h-3v-11h3v11zm-1.5-12.268c-.966 0-1.
75-.79-1.75-1.764s.784-1.764 1.75-1.764 1.75.79 1.75 1.764-.783 1.764-1.75 1.764zm13.5 12.268h-3v-5.604c0-3.368-4-
3.113-4 0v5.604h-3v-11h3v1.765c1.396-2.586 7-2.777 7 2.476v6.759z"/></svg>
278     </a>
279
280     <div style="flex: 1; padding: 0 15px; position: relative; display: flex; align-items: center; justify-conten
281     <div style="position: absolute; left: 0; width: 100%; text-align: center; font-family: &#x27;Space Grotes
k', sans-serif; font-size: 0.9rem; color: var(--accent); letter-spacing: 0.5px; z-index: 1;">This application was d
eveloped as a coursework assignment for the Information Retrieval module, MSc in Data Science & Computational I
ntelligence, Coventry University.</div>
282     <div style="position: relative; z-index: 2; padding: 0 10px; background: var(--bg-0); box-shadow: 0 0 15px
15px var(--bg-0); border-radius: 10px;">
283         <span class="gradient-text" style="font-family: &#x27;Dancing Script', cursive; font-size: 0.85rem; lett
er-spacing: 1px;">Developed By Rohit Jha Coventry University _ 11782276</span>
284     </div>
285 </div>
286
287     <a href="https://github.com/jrohitofficial" target="_blank" title="GitHub" style="color: var(--muted); displ
ay: flex; transition: all 0.2s ease;">
288     <svg width="18" height="18" viewBox="0 0 24 24" fill="currentColor"><path d="M12 0c-6.626 0-12 5.373-12 12
0 5.302 3.438 9.8 8.207 11.387.599.111.793-.261.793-.577v-2.234c-3.338-.726-4.033-1.416-.546-1.387-1.333
-1.756-1.333-1.756-1.089-.745.083-.729.083-.729 1.205.084 1.839 1.237 1.839 1.237 1.07 1.834 2.807 1.304 3.492.997.
107-.775.418-1.305.762-1.604-2.665-.305-5.467-1.334-5.467-5.931 0-1.311.469-2.381 1.236-3.221-.124-.303-.535-1.524.
117-3.176 0 0 1.008-.322 3.301 1.23.957-.266 1.983-.399 3.003-.404 1.02.005 2.047.138 3.006.404 2.291-1.552 3.297-
1.23 3.297-1.23.653 1.653.242 2.874.118 3.176.77.84 1.235 1.911 1.235 3.221 0 4.609-2.807 5.624-5.479 5.921.43.372.
823 1.102.823 2.222v3.293c0 .319.192.694.801.576 4.765-1.589 8.199-6.086 8.199-11.386 0-6.627-5.373-12-12z"/></s
vg>
289     </a>
290 </div>
291 </footer>
292
293 <script>
294     window.TASK1_BASE = "{{ task1_url }}";
295     window.TASK2_BASE = "{{ task2_url }}";
296 </script>
297 <script src="{{ url_for(&#x27;static', filename='js/app.js') }}"?v=17"></script>
298 </body>
299
300 </html>
```

Source Code: [frontend_html_part6](#)

```
app.js (Part 1)

1  const T1 = window.TASK1_BASE;
2  const T2 = window.TASK2_BASE;
3
4  /* ----- Tab switching ----- */
5  const tabBtns = document.querySelectorAll(".tab-btn");
6  const indicator = document.getElementById("tab-indicator");
7
8  function positionIndicator(btn) {
9      indicator.style.width = btn.offsetWidth + "px";
10     indicator.style.transform = `translateX(${btn.offsetLeft - 4}px)`;
11 }
12
13 function activateTab(name) {
14     tabBtns.forEach((b) => {
15         const active = b.dataset.tab === name;
16         b.classList.toggle("active", active);
17         b.setAttribute("aria-selected", active);
18         if (active) positionIndicator(b);
19     });
20     document.getElementById("panel-search").classList.toggle("active", name === "search");
21     document.getElementById("panel-cluster").classList.toggle("active", name === "cluster");
22 }
23
24 tabBtns.forEach((b) => b.addEventListener("click", () => activateTab(b.dataset.tab)));
25 window.addEventListener("load", () => positionIndicator(document.querySelector(".tab-btn.active")));
26 window.addEventListener("resize", () => positionIndicator(document.querySelector(".tab-btn.active")));
27
28 /* ----- Backend status pills ----- */
29 async function checkStatus(url, pillId) {
30     const pill = document.getElementById(pillId);
31     try {
32         const res = await fetch(url, { method: "GET" });
33         pill.classList.add(res.ok ? "online" : "offline");
34     } catch {
35         pill.classList.add("offline");
36     }
37 }
38 checkStatus(`${T1}/api/crawler/status`, "status-task1");
39 checkStatus(`${T2}/api/dataset/stats`, "status-task2");
40
41 /* ----- Subtle 3D tilt on cards ----- */
42 function attachTilt(el) {
43     el.addEventListener("mousemove", (e) => {
44         const r = el.getBoundingClientRect();
45         const px = (e.clientX - r.left) / r.width - 0.5;
46         const py = (e.clientY - r.top) / r.height - 0.5;
47         el.style.transform = `perspective(800px) rotateX(${(-py * 4).toFixed(2)}deg) rotateY(${(px * 4).toFixed(2)}deg) translateY(-2px)`;
48     });
49     el.addEventListener("mouseleave", () => { el.style.transform = ""; });
50 }
51
52 /* ===== TASK 1: SEARCH ===== */
53 const searchForm = document.getElementById("search-form");
54 const searchInput = document.getElementById("search-input");
55 const resultsEl = document.getElementById("results");
```

Source Code: frontend_js_part1


```

app.js (Part 10)

496   if (query.length < 2) {
497       task1Suggestions.classList.add("hidden");
498       return;
499   }
500   try {
501       const res = await fetch(`${T1}/api/suggest?q=${encodeURIComponent(query)}`);
502       const data = await res.json();
503       if (data.length > 0) {
504           task1Suggestions.innerHTML = data.map(s => `<li>${escapeHtml(s)}</li>`).join("");
505           task1Suggestions.classList.remove("hidden");
506       } else {
507           task1Suggestions.classList.add("hidden");
508       }
509   } catch {
510       task1Suggestions.classList.add("hidden");
511   }
512 }, 300));
513
514 task1Suggestions.addEventListener("click", (e) => {
515     if (e.target.tagName === "LI") {
516         searchInput.value = e.target.textContent;
517         task1Suggestions.classList.add("hidden");
518         searchForm.dispatchEvent(new Event("submit", { cancelable: true, bubbles: true }));
519     }
520 });
521
522 // Task 2 Suggest
523 const task2Suggestions = document.getElementById("task2-suggestions");
524 if (task2Suggestions) {
525     textInput.addEventListener("input", debounce(async (e) => {
526         const query = e.target.value.trim();
527         if (query.length < 2) {
528             task2Suggestions.classList.add("hidden");
529             return;
530         }
531         try {
532             const res = await fetch(`${T2}/api/suggest?q=${encodeURIComponent(query)}`);
533             const data = await res.json();
534             if (data.length > 0) {
535                 let html = `<li class="close-suggestions" style="text-align: right; padding: 6px 12px; color: var(--rose);
cursor: pointer; font-size: 0.85rem; border-bottom: 1px solid rgba(255,255,255,0.05); display: flex; justify-content:
t: flex-end; align-items: center; gap: 4px;"><svg viewBox="0 0 24 24" fill="none" stroke="currentColor" stroke-wid
th="2" style="width: 14px; height: 14px;"><line x1="18" y1="6" x2="6" y2="18"></line><line x1="6" y1="6" x2="18" y
2="18"></line></svg> Close</li>`;
536                 html += data.map(s => `<li>${escapeHtml(s)}</li>`).join("");
537                 task2Suggestions.innerHTML = html;
538                 task2Suggestions.classList.remove("hidden");
539             } else {
540                 task2Suggestions.classList.add("hidden");
541             }
542         } catch {
543             task2Suggestions.classList.add("hidden");
544         }
545     }, 300));
546
547     task2Suggestions.addEventListener("click", (e) => {
548         const li = e.target.closest("li");
549         if (!li) return;
550     });

```

Source Code: [frontend_js_part10](#)

```
app.js (Part 11)

551     if (li.classList.contains("close-suggestions")) {
552         task2Suggestions.classList.add("hidden");
553         return;
554     }
555
556     textInput.value = li.textContent;
557     task2Suggestions.classList.add("hidden");
558 });
559 }
560
561 const retrainBtn = document.getElementById("retrain-btn");
562 if (retrainBtn) {
563     retrainBtn.addEventListener("click", async () => {
564         retrainBtn.disabled = true;
565         retrainBtn.innerHTML = `<span class="spinner" style="width:14px;height:14px;border:2px solid var(--accent);border-bottom-color:transparent;border-radius:50%;animation:spin 1s linear infinite;display:inline-block;"></span> Collecting Data...`;
566         showToast("Started data collection. This takes ~60s due to remote DB.");
567         try {
568             await fetch(`${T2}/api/retrain`, { method: "POST" });
569             setTimeout(() => {
570                 showToast("Training new K-Means model...");
571                 retrainBtn.innerHTML = `<span class="spinner" style="width:14px;height:14px;border:2px solid var(--accent);border-bottom-color:transparent;border-radius:50%;animation:spin 1s linear infinite;display:inline-block;"></span> Training Model...`;
572                 setTimeout(() => {
573                     loadDatasetStats();
574                     loadModelStats();
575                     loadPCAChart();
576                     retrainBtn.disabled = false;
577                     retrainBtn.innerHTML = `<svg viewBox="0 0 24 24" fill="none" style="width:14px;height:14px;"><path d="M4 4v5h.582m15.356 2A8.001 8.001 0 004.582 9m0 0H9m11 11v-5h-.581m0 0a8.003 8.003 0 01-15.357-2m15.357 2H15" stroke="currentColor" stroke-width="2" stroke-linecap="round" stroke-linejoin="round"/></svg> Collect & Retrain`;
578                     showToast("Retraining complete! Stats updated.");
579                     }, 30000); // 30 seconds for training phase
580                 }, 35000); // 35 seconds for collection phase
581             } catch (err) {
582                 showToast("Failed to start retraining.");
583                 retrainBtn.disabled = false;
584                 retrainBtn.innerHTML = "Collect & Retrain";
585             }
586         });
587     }
588
589     const refreshHistoryBtn = document.getElementById("refresh-history-btn");
590     if (refreshHistoryBtn) {
591         refreshHistoryBtn.addEventListener("click", () => {
592             const icon = refreshHistoryBtn.querySelector("svg");
593             if (icon) {
594                 icon.style.animation = "spin 0.5s linear";
595                 setTimeout(() => icon.style.animation = "", 500);
596             }
597             loadHistory();
598         });
599     }
600
601     const clearHistoryBtn = document.getElementById("clear-history-btn");
602     const confirmModal = document.getElementById("confirm-modal");
603     const confirmYesBtn = document.getElementById("confirm-yes-btn");
604     const confirmCancelBtn = document.getElementById("confirm-cancel-btn");
605 }
```

Source Code: [frontend_js_part11](#)

```
app.js (Part 12)

606 if (clearHistoryBtn && confirmModal && confirmYesBtn && confirmCancelBtn) {
607   clearHistoryBtn.addEventListener("click", () => {
608     confirmModal.classList.remove("hidden");
609   });
610
611   confirmCancelBtn.addEventListener("click", () => {
612     confirmModal.classList.add("hidden");
613   });
614
615   confirmYesBtn.addEventListener("click", async () => {
616     confirmModal.classList.add("hidden");
617     try {
618       const res = await fetch(`${T2}/api/predictions/history`, { method: "DELETE" });
619       if (res.ok) {
620         loadHistory();
621         showToast("History cleared.");
622       } else {
623         throw new Error();
624       }
625     } catch (err) {
626       showToast("Failed to clear history.");
627     }
628   });
629 }
630
631 // Hide on outside click
632 document.addEventListener("click", (e) => {
633   if (!e.target.closest("#search-form") && !e.target.closest("#task1-suggestions")) {
634     task1Suggestions.classList.add("hidden");
635   }
636   if (!e.target.closest(".classify-card") && !e.target.closest("#task2-suggestions") && task2Suggestions) {
637     task2Suggestions.classList.add("hidden");
638   }
639 });
640
641 // Help Modal Logic
642 const helpBtn = document.getElementById("help-btn");
643 const helpModal = document.getElementById("help-modal");
644 const closeHelpBtn = document.getElementById("close-help-btn");
645
646 if (helpBtn && helpModal && closeHelpBtn) {
647   helpBtn.addEventListener("click", () => {
648     helpModal.classList.remove("hidden");
649   });
650
651   closeHelpBtn.addEventListener("click", () => {
652     helpModal.classList.add("hidden");
653   });
654
655   helpModal.addEventListener("click", (e) => {
656     if (e.target === helpModal) {
657       helpModal.classList.add("hidden");
658     }
659   });
660 }
```

Source Code: [frontend_js_part12](#)

```
app.js (Part 2)

56 const metaEl = document.getElementById("results-meta");
57 const paginationEl = document.getElementById("pagination");
58 const statusToggle = document.getElementById("crawler-status-toggle");
59 const statusPanel = document.getElementById("crawler-status-panel");
60
61 let currentQuery = "";
62 let currentPage = 1;
63
64 searchForm.addEventListener("submit", (e) => {
65   e.preventDefault();
66   currentQuery = searchInput.value.trim();
67   currentPage = 1;
68   if (currentQuery) runSearch();
69 });
70
71 document.querySelectorAll(".chip[data-q]").forEach((chip) => {
72   chip.addEventListener("click", () => {
73     searchInput.value = chip.dataset.q;
74     currentQuery = chip.dataset.q;
75     currentPage = 1;
76     runSearch();
77   });
78 });
79
80 let statusInterval = null;
81
82 async function fetchStatus(isInitial = false) {
83   if (statusPanel.classList.contains("hidden")) return;
84   if (isInitial) statusPanel.innerHTML = `<p class="state-msg">Loading crawler status...</p>`;
85   try {
86     const res = await fetch(`${T1}/api/crawler/status`);
87     const data = await res.json();
88     let lastCrawl = "No crawl has run yet.";
89     if (data.last_crawl) {
90       if (data.last_crawl.stopped_reason === "Running...") {
91         lastCrawl = `<span style="color: var(--rose); font-weight: 600; display: inline-flex; align-items: center; gap: 8px;"><span class="spinner" style="width: 14px; height: 14px; border: 2px solid var(--rose); border-bottom-color: transparent; border-radius: 50%; animation: spin 1s linear infinite;"></span> Crawling live right now... (${data.research_output_count} pubs & ${data.profile_count} profiles found so far!)</span>`;
92       } else {
93         let finishedDate = data.last_crawl.finished_at ? new Date(data.last_crawl.finished_at).toLocaleString() : `<span>Unknown date</span>`;
94         let remainingDaysText = `<span>`;
95         if (data.scheduler && data.scheduler.jobs && data.scheduler.jobs.length > 0 && data.scheduler.jobs[0].next_run_time) {
96           let nextRun = new Date(data.scheduler.jobs[0].next_run_time);
97           let diffMs = nextRun - new Date();
98           let days = Math.ceil(diffMs / (1000 * 60 * 60 * 24));
99           remainingDaysText = ` (Next auto-crawl in ${days} days)`;
100         }
101         lastCrawl = `<span style="color: var(--text); font-weight: 500;">Last Crawled: ${finishedDate}${remainingDaysText}</span>`;
102       }
103     }
104
105     if (!document.getElementById("spin-style")) {
106       document.head.insertAdjacentHTML("beforeend", `<style id="spin-style">@keyframes spin { to { transform: rotate(360deg); } }</style>`);
107     }
108
109     statusPanel.innerHTML = `
110     <table>
```

Source Code: [frontend_js_part2](#)

```

111         <tr><td>Research outputs indexed</td><td>${data.research_output_count}</td></tr>
112         <tr><td>Profiles indexed</td><td>${data.profile_count}</td></tr>
113         <tr><td>TF-IDF Search Index</td><td>
114             <span style="color: ${data.index_ready ? '#x27;var(--success)' : 'var(--accent)'}; font-weight: bold;">
115                 ${data.index_ready ? "Ready & Optimized" : "Building..."}
116             </span>
117         </td></tr>
118         <tr><td>Auto-Crawl Scheduler</td><td>
119             <span style="color: var(--success); font-weight: bold;">${data.scheduler.running ? "Active" : "Inactive"}</span>
120             <span style="color: var(--muted); font-size: 0.9em;">(Runs every ${data.scheduler.interval_months ?? "?"}
months)</span>
121         </td></tr>
122         <tr><td>Last crawl</td><td>${lastCrawl}</td></tr>
123     </table>`;
124     } catch {
125         statusPanel.innerHTML = `&#x27;<p class="state-msg error">Failed to reach the Task 1 backend on :5001.</p>`;
126     }
127 }
128
129 statusToggle.addEventListener("click", () => {
130     document.getElementById("search-center-wrapper").style.marginTop = "0";
131     statusPanel.classList.toggle("hidden");
132     if (!statusPanel.classList.contains("hidden")) {
133         fetchStatus(true);
134         if (!statusInterval) statusInterval = setInterval(() => fetchStatus(false), 3000);
135     } else {
136         if (statusInterval) {
137             clearInterval(statusInterval);
138             statusInterval = null;
139         }
140     }
141 });
142
143 async function runSearch() {
144     document.getElementById("search-center-wrapper").style.marginTop = "0";
145     resultsEl.innerHTML = `&#x27;<p class="state-msg">Searching...</p>`;
146     metaEl.classList.add("hidden");
147     paginationEl.classList.add("hidden");
148     try {
149         const res = await fetch(`${T1}/api/search?q=${encodeURIComponent(currentQuery)}&page=${currentPage}`);
150         if (!res.ok) throw new Error((await res.json().catch(() => ({}))).error || `Request failed (${res.status})`);
151         renderResults(await res.json());
152     } catch (err) {
153         resultsEl.innerHTML = `<p class="state-msg error">${err.message}. Is the Task 1 backend running on :5001?</p>`;
154     }
155 }
156
157 function renderResults(data) {
158     metaEl.classList.remove("hidden");
159     if (data.total_results === 0) {
160         metaEl.textContent = `No results for "${data.query}"`;
161         resultsEl.innerHTML = `&#x27;<p class="state-msg">Try a different keyword or author name.</p>`;
162         return;
163     }
164     metaEl.textContent = `${data.total_results} result${data.total_results === 1 ? "" : "s"} for "${data.query}" | p
age ${data.page} of ${data.total_pages}`;
165     resultsEl.innerHTML = data.results.map(renderCard).join("");

```

Source Code: [frontend_js_part3](#)

```

166     resultsEl.querySelectorAll(".result-card").forEach(attachTilt);
167     renderPagination(data);
168 }
169
170 function renderCard(r, i) {
171     const authorsHtml = r.authors
172       .map((name, idx) => {
173         const url = r.author_profiles[idx];
174         return url ? `<a href="${url}" target="_blank" rel="noopener">${escapeHtml(name)}</a>` : escapeHtml(name);
175       })
176       .join(" · ");
177     return `
178       <article class="result-card">
179         <span class="result-rank">#${(currentPage - 1) * 10 + i + 1}</span>
180         <div class="result-title"><a href="${r.document_url}" target="_blank" rel="noopener">${escapeHtml(r.title)}</a></div>
181         <div class="result-meta"><span class="authors">${authorsHtml || "Authors not listed"}</span>${r.publication_type ? ` · ${escapeHtml(r.publication_type)}` : ""}${r.journal ? ` · ${escapeHtml(r.journal)}` : ""}</div>
182         <div class="result-desc">${escapeHtml(truncate(r.description, 240))}</p>` : ""}
183         <div class="result-footer">
184           <span style="color:var(--muted);font-size:.8rem;">Published: ${escapeHtml(r.publication_date || "Not state
185             d")}</span>
186           <span class="similarity-badge">Cosine similarity: ${r.cosine_similarity.toFixed(4)}</span>
187           <a class="view-link" href="${r.document_url}" target="_blank" rel="noopener">View Research Output </a>
188         </div>
189       </article>`;
190 }
191
192 function renderPagination(data) {
193     if (data.total_pages <= 1) { paginationEl.classList.add("hidden"); return; }
194     paginationEl.classList.remove("hidden");
195     let html = `<button ${data.page <= 1 ? "disabled" : ""} data-page="${data.page - 1}">Previous</button>`;
196     for (let p = 1; p <= data.total_pages; p++) html += `<button class="${p === data.page ? "active" : ""}" data-page="${p}">${p}</button>`;
197     html += `<button ${data.page >= data.total_pages ? "disabled" : ""} data-page="${data.page + 1}">Next</button>`;
198     paginationEl.innerHTML = html;
199     paginationEl.querySelectorAll("button[data-page]").forEach((btn) => {
200       btn.addEventListener("click", () => { currentPage = parseInt(btn.dataset.page, 10); runSearch(); window.scrollTo({ top: 0, behavior: "smooth" }); });
201     });
202 }
203
204 /* ===== TASK 2: CLUSTERING ===== */
205 const textInput = document.getElementById("text-input");
206 const classifyBtn = document.getElementById("classify-btn");
207 const clearBtn = document.getElementById("clear-btn");
208 const resultPanel = document.getElementById("result-panel");
209 const historyList = document.getElementById("history-list");
210
211 classifyBtn.addEventListener("click", classify);
212 clearBtn.addEventListener("click", () => { textInput.value = ""; resultPanel.classList.add("hidden"); textInput.rows = 7; });
213
214 const CAT_COLORS = { Economics: "#22d3c9", Entertainment: "#f5a524", Politics: "#a78bfa" };
215
216 function showToast(msg) {
217     const toast = document.createElement("div");
218     toast.className = "toast-popup";
219     toast.textContent = msg;
220     document.body.appendChild(toast);
221     setTimeout(() => toast.classList.add("show"), 10);
222 }

```

Source Code: [frontend_js_part4](#)

```
app.js (Part 5)

221   setTimeout(() => {
222     toast.classList.remove("show");
223     setTimeout(() => toast.remove(), 300);
224   }, 2500);
225 }
226
227 async function classify() {
228   const text = textInput.value.trim();
229   const loader = document.getElementById("classify-loader");
230
231   if (!text) {
232     resultPanel.classList.remove("hidden");
233     resultPanel.innerHTML = `&#x27;<p class="state-msg error">Please enter some text to classify.</p>`;
234     return;
235   }
236
237   loader.classList.remove("hidden");
238   resultPanel.classList.add("hidden");
239
240   try {
241     const res = await fetch(`${T2}/api/classify`, { method: "POST", headers: { "Content-Type": "application/json"
242 }, body: JSON.stringify({ text }) });
243     const data = await res.json().catch(() => ({}));
244
245     loader.classList.add("hidden");
246
247     if (!res.ok) {
248       if (res.status === 400 && data.error) {
249         showToast(data.error);
250         return; // do not show error inside the panel
251       }
252       throw new Error(data.error || "Classification failed.");
253     }
254
255     resultPanel.classList.remove("hidden");
256     textInput.rows = 4;
257     const T = 0.04;
258     let sumExp = 0;
259     const exps = {};
260     for (const [cat, d] of Object.entries(data.distances_all_clusters)) {
261       const e = Math.exp(-d / T);
262       exps[cat] = e;
263       sumExp += e;
264     }
265     let confidence = Math.round((exps[data.predicted_category] / sumExp) * 100);
266     if (isNaN(confidence)) confidence = 99; // fallback if math overflows
267     const color = CAT_COLORS[data.predicted_category] || "#60a5fa";
268
269     let iconSvg = "";
270     if (data.predicted_category === "Politics") {
271       iconSvg = `<svg width="24" height="24" viewBox="0 0 24 24" fill="none" stroke="${color}" stroke-width="2" st
272 roke-linecap="round" stroke-linejoin="round"><rect x="4" y="10" width="16" height="10"></rect><path d="M12 2L3 10h
273 18z"></path><path d="M8 10v10"></path><path d="M16 10v10"></path></svg>`;
274     } else if (data.predicted_category === "Economics") {
275       iconSvg = `<svg width="24" height="24" viewBox="0 0 24 24" fill="none" stroke="${color}" stroke-width="2" st
276 roke-linecap="round" stroke-linejoin="round"><line x1="12" y1="1" x2="12" y2="23"></line><path d="M17 5H9.5a3.5 3.
277 5 0 0 0 7h5a3.5 3.5 0 0 1 0 7H6"></path></svg>`;
278     } else {
279       iconSvg = `<svg width="24" height="24" viewBox="0 0 24 24" fill="none" stroke="${color}" stroke-width="2" st
280 roke-linecap="round" stroke-linejoin="round"><polygon points="5 3 19 12 5 21 5 3"></polygon></svg>`;
281     }
282   }
283 }
```

Source Code: frontend_js_part5

```

276
277     resultPanel.innerHTML = `
278         <div style="text-align: center; padding: 0;">
279             <div class="result-category ${data.predicted_category}" style="line-height: 1; padding-bottom: 8px;">${dat
280                 <div style="display: flex; align-items: center; justify-content: center; gap: 10px; font-size: 0.85rem; color:
281                 var(--muted); margin-top: 0;">
282                     <span style="white-space: nowrap;">Confidence: ${confidence}%</span>
283                     <div style="width: 120px; height: 6px; background: rgba(255,255,255,0.1); border-radius: 3px; overflow: h
284                     idden;">
285                         <div style="height: 100%; width: ${confidence}%; background: ${color}; transition: width 1s ease-out;">
286                     </div>
287                 </div>
288             </div>
289         </div>
290     `;
291     loadHistory();
292 } catch (err) {
293     loader.classList.add("hidden");
294     resultPanel.classList.remove("hidden");
295     textInput.rows = 7;
296     resultPanel.innerHTML = `<p class="state-msg error">${err.message}. Is the Task 2 backend running on :5002?</p>`;
297 }
298 }
299
300 async function loadHistory() {
301     try {
302         const res = await fetch(`${T2}/api/predictions/history?limit=8&t=${Date.now()}`);
303         const data = await res.json();
304         if (!data.predictions.length) { historyList.innerHTML = `<p class="state-msg">No predictions yet.</p>`; re
305         turn; }
306         historyList.innerHTML = data.predictions.map((p) => `
307             <div class="history-item">
308                 <span class="cat-tag ${p.predicted_category}">${p.predicted_category}</span>
309                 <span class="history-text">${escapeHtml(p.input_text)}</span>
310                 <span class="history-ts">${new Date(p.timestamp).toLocaleString()}</span>
311             </div>`);
312         } catch {
313             historyList.innerHTML = `<p class="state-msg error">Failed to load history.</p>`;
314         }
315     }
316 }
317
318 /* ----- Dataset / Model stats tabs ----- */
319 const statsTabs = document.querySelectorAll(".stats-tab");
320 const datasetBody = document.getElementById("dataset-stats");
321 const modelBody = document.getElementById("model-stats");
322
323 statsTabs.forEach((tab) => tab.addEventListener("click", () => {
324     statsTabs.forEach((t) => t.classList.remove("active"));
325     tab.classList.add("active");
326     datasetBody.classList.toggle("hidden", tab.dataset.stats !== "dataset");
327     modelBody.classList.toggle("hidden", tab.dataset.stats !== "model");
328 }));
329
330 async function loadDatasetStats() {
331     try {
332         const res = await fetch(`${T2}/api/dataset/stats`);
333         const data = await res.json();
334
335         let currentAngle = 0;

```

Source Code: [frontend_js_part6](#)


```

331     const radius = 60;
332     const circumference = 2 * Math.PI * radius;
333     const total = data.total;
334
335     const colors = { Economics: "var(--econ)", Entertainment: "var(--ent)", Politics: "var(--pol)" };
336     const icons = { Economics: "📊", Entertainment: "🎮", Politics: "🏛️" };
337
338     let svg = `<div style="perspective: 800px; width:180px; height:180px; margin:0 auto;"><svg viewBox="0 0 160 160" style="width:100%; height:100%; display:block; filter: drop-shadow(0 25px 20px rgba(0,0,0,0.7)) drop-shadow(0 0 10px rgba(255,255,255,0.1)); transform: rotateX(25deg) rotateY(-15deg) scale(1.1); transform-style: preserve-3d; transition: transform 0.5s;">`;
339     let legend = `<div style="display:flex; flex-direction:column; gap:12px;">`;
340
341     data.categories.forEach((c) => {
342         const cat = c.category;
343         const count = c.actual;
344         const percent = count / total;
345         const strokeDasharray = `${percent * circumference} ${circumference}`;
346         const strokeDashoffset = - (currentAngle / 360) * circumference;
347
348         svg += `<circle r="${radius}" cx="80" cy="80" fill="transparent" stroke="${colors[cat]}" stroke-width="24" stroke-dasharray="${strokeDasharray}" stroke-dashoffset="${strokeDashoffset}" transform="rotate(-90 80 80)" style="transition: all 1s ease-out; stroke-linecap: round; filter: drop-shadow(inset 0 4px 4px rgba(255,255,255,0.5));" />`;
349         currentAngle += percent * 360;
350
351         legend += `<div style="display:flex; align-items:center; gap:14px; font-size:14px;">
352             <div style="width:12px; height:12px; border-radius:50%; background:${colors[cat]}; box-shadow: 0 0 12px ${colors[cat]};"></div>
353             <div style="color:var(--muted); display:flex; align-items:center; gap:8px;">
354                 <span>${icons[cat]}</span>
355                 <span>${cat}</span>
356             </div>
357             <div class="count-up" data-count="${count}" style="color:#fff; font-weight:800; font-size:15px; margin-left:auto;">0</div>
358             </div>`;
359     });
360
361     svg += `<text x="80" y="75" class="count-up" data-count="${total}" text-anchor="middle" fill="#fff" font-size="28" font-weight="900" font-family="var(--font-sans)" style="transform: translateZ(20px); text-shadow: 0 5px 10px rgba(0,0,0,0.8);">0</text>
362         <text x="80" y="98" text-anchor="middle" fill="var(--muted)" font-size="12" font-family="var(--font-sans)" style="transform: translateZ(20px);">documents</text>
363     </svg></div>`;
364     legend += `</div>`;
365
366     const html = `
367     <div style="margin-bottom: 8px;">
368         <h3 style="color:#fff; font-size:1.05rem; font-weight:800; margin-bottom: 16px; text-align:center;">Cluster Distribution</h3>
369         <div style="display:flex; justify-content:center; align-items:center; gap: 32px;">
370             ${svg}
371             ${legend}
372         </div>
373     </div>`;
374     datasetBody.innerHTML = html;
375
376     // Animate numbers
377     datasetBody.querySelectorAll(`&#x27;.count-up`).forEach(el => {
378         const target = parseInt(el.getAttribute(`&#x27;data-count`), 10);
379         const duration = 1500;
380         const startTime = performance.now();
381
382         function updateNumber(currentTime) {
383             const elapsed = currentTime - startTime;
384             const progress = Math.min(elapsed / duration, 1);
385             const easeOutQuad = progress * (2 - progress);

```

Source Code: [frontend_js_part7](#)

```

386         el.textContent = Math.floor(easeOutQuad * target);
387         if (progress < 1) requestAnimationFrame(updateNumber);
388         else el.textContent = target;
389     }
390     requestAnimationFrame(updateNumber);
391 });
392 } catch (err) {
393     datasetBody.innerHTML = `<p class="state-msg error">Failed to reach the Task 2 backend on :5002.</p>`;
394 }
395 }
396
397 async function loadModelStats() {
398     try {
399         const res = await fetch(`${T2}/api/model/evaluation`);
400         if (!res.ok) { modelBody.innerHTML = `<p class="state-msg">Model not trained yet. Run scripts/train_model.py.</p>`; return; }
401         const d = await res.json();
402         modelBody.innerHTML = `
403         <table>
404             <tr><td>Documents / vocabulary terms</td><td>${d.n_documents} / ${d.n_terms}</td></tr>
405             <tr><td>Silhouette score</td><td>${d.silhouette_score.toFixed(4)}</td></tr>
406             <tr><td>Inertia</td><td>${d.inertia.toFixed(2)}</td></tr>
407             <tr><td>Accuracy vs. known labels</td><td>${(d.accuracy * 100).toFixed(1)}%</td></tr>
408             <tr><td>Precision (macro)</td><td>${d.precision_macro.toFixed(4)}</td></tr>
409             <tr><td>Recall (macro)</td><td>${d.recall_macro.toFixed(4)}</td></tr>
410             <tr><td>F1 (macro)</td><td>${d.f1_macro.toFixed(4)}</td></tr>
411         </table>`;
412     } catch {
413         modelBody.innerHTML = `<p class="state-msg error">Failed to reach the Task 2 backend on :5002.</p>`;
414     }
415 }
416
417 let currentPCAChart = null;
418
419 async function loadPCAChart() {
420     const ctx = document.getElementById(`<x27;clusterCanvas`);
421     if (!ctx) return;
422
423     try {
424         const res = await fetch(`${T2}/api/model/pca`);
425         if (!res.ok) throw new Error("Not trained");
426         const data = await res.json();
427
428         const colors = { Economics: "rgba(34, 211, 201, 0.8)", Entertainment: "rgba(245, 165, 36, 0.8)", Politics: "rg
429 ba(167, 139, 250, 0.8)" };
430         const borders = { Economics: "#22d3c9", Entertainment: "#f5a524", Politics: "#a78bfa" };
431
432         const datasets = ["Economics", "Entertainment", "Politics"].map(cat => {
433             return {
434                 label: cat,
435                 data: data.filter(d => d.category === cat).map(d => ({ x: d.x, y: d.y, title: d.title })),
436                 backgroundColor: colors[cat],
437                 borderColor: borders[cat],
438                 borderWidth: 1,
439                 pointRadius: 4,
440                 pointHoverRadius: 7
441             };

```

Source Code: [frontend_js_part8](#)

```

441     });
442
443     if (currentPCAChart) {
444         currentPCAChart.destroy();
445     }
446
447     currentPCAChart = new Chart(ctx, {
448         type: 'scatter',
449         data: { datasets },
450         options: {
451             responsive: true,
452             maintainAspectRatio: false,
453             plugins: {
454                 legend: { display: false },
455                 tooltip: {
456                     callbacks: {
457                         label: function(context) {
458                             return context.raw.title;
459                         }
460                     }
461                 },
462             },
463             scales: {
464                 x: { grid: { color: "#2a2d3e" }, ticks: { display: false } },
465                 y: { grid: { color: "#2a2d3e" }, ticks: { display: false } }
466             }
467         }
468     });
469 } catch (err) {
470     const container = ctx.parentNode;
471     container.innerHTML = `<div style="padding: 20px; text-align: center; color: var(--muted);">Error loading PCA chart: ${err.message}. Run scripts/train_model.py to generate PCA data.</div>`;
472 }
473 }
474 loadPCAChart();
475
476 function truncate(text, n) { return text.length > n ? text.slice(0, n).trim() + "..." : text; }
477 function escapeHtml(str) { const d = document.createElement("div"); d.textContent = str; return d.innerHTML; }
478
479 loadHistory();
480 loadDatasetStats();
481 loadModelStats();
482
483 /* ===== AUTO-SUGGEST ===== */
484 function debounce(func, wait) {
485     let timeout;
486     return function(...args) {
487         clearTimeout(timeout);
488         timeout = setTimeout(() => func.apply(this, args), wait);
489     };
490 }
491
492 // Task 1 Suggest
493 const task1Suggestions = document.getElementById("task1-suggestions");
494 searchInput.addEventListener("input", debounce(async (e) => {
495     const query = e.target.value.trim();

```

Source Code: frontend_js_part9

```
settings.py (Part 1)

1  """
2  Central configuration for the Task 1 vertical search engine backend.
3
4  All values are read from environment variables (see .env.example at the
5  project root). Nothing sensitive is hard-coded here.
6  """
7  import os
8  from pathlib import Path
9  from dotenv import load_dotenv
10
11  # Load the .env file that sits at the IR_COURSEWORK project root, two levels
12  # above this file (backend/config/settings.py -> backend -> task1_vertical_search -> IR_COURSEWORK).
13  PROJECT_ROOT = Path(__file__).resolve().parents[3]
14  load_dotenv(PROJECT_ROOT / ".env")
15
16
17  class Settings:
18      # --- MongoDB ---
19      MONGODB_URI: str = os.environ.get("MONGODB_URI", "mongodb://localhost:27017")
20      DATABASE_NAME: str = os.environ.get("TASK1_DATABASE_NAME", "Task_1_Vertical_Search")
21
22      # --- Crawl target ---
23      # The exact seed URL mandated by the ST7071CEM coursework brief. Do not
24      # change this to a different organisation.
25      SEED_URL: str = os.environ.get(
26          "SEED_URL",
27          "https://pureportal.coventry.ac.uk/en/organisations/"
28          "centre-for-healthcare-and-community-transformation/",
29      )
30      CRAWLER_USER_AGENT: str = os.environ.get(
31          "CRAWLER_USER_AGENT",
32          "IRCOURSEWORKBot/1.0 (Coventry ST7071CEM coursework crawler)",
33      )
34      # robots.txt for pureportal.coventry.ac.uk publishes "Crawl-delay: 5".
35      # This value is read at runtime rather than hard-coded so it can be
36      # raised if the site's policy changes.
37      CRAWL_DELAY_SECONDS: float = float(os.environ.get("CRAWL_DELAY_SECONDS", "5"))
38
39      # Breadth-first crawl bounds so the crawler provably cannot run forever
40      # (coursework requirement: "must not endlessly crawl the Coventry website").
41      MAX_PROFILES: int = int(os.environ.get("MAX_PROFILES", "200"))
42      MAX_PUBLICATIONS: int = int(os.environ.get("MAX_PUBLICATIONS", "200"))
43      MAX_CRAWL_SECONDS: int = int(os.environ.get("MAX_CRAWL_SECONDS", "1800"))
44
45      # --- Scheduler ---
46      # The interval is kept as an explicit, documented, *configurable* value.
47      # The coursework now requires a 3-month default schedule.
48      CRAWL_INTERVAL_MONTHS: int = int(os.environ.get("CRAWL_INTERVAL_MONTHS", "3"))
49      RUN_CRAWL_ON_STARTUP: bool = os.environ.get("RUN_CRAWL_ON_STARTUP", "false").lower() == "true"
50
51      # --- API / ranking ---
52      TOP_K: int = int(os.environ.get("TOP_K", "10"))
53      TASK1_PORT: int = int(os.environ.get("TASK1_PORT", "5001"))
54
55      LOG_DIR: Path = Path(__file__).resolve().parents[1] / "logs"
```

Source Code: [task1_config_settings_part1](#)

```
settings.py (Part 2)

56
57
58  settings = Settings()
59  settings.LOG_DIR.mkdir(parents=True, exist_ok=True)
60
```

Source Code: [task1_config_settings_part2](#)

```
http_client.py (Part 1)

1  """
2  Polite HTTP fetching for the Task 1 crawler.
3
4  Responsibilities:
5      - honour robots.txt (allow-check + crawl-delay) before every request
6      - identify the crawler with a descriptive User-Agent
7      - apply request timeouts and a bounded number of retries with backoff
8      - rate-limit consecutive requests to the same host
9  """
10 import logging
11 import time
12 from urllib.parse import urlparse
13
14 import requests
15
16 from config.settings import settings
17 from crawler.robots_check import crawl_delay, is_allowed
18
19 logger = logging.getLogger("task1.crawler.http")
20
21 _last_request_time: dict[str, float] = {}
22
23
24 class FetchError(Exception):
25     pass
26
27
28 class RobotsDisallowed(FetchError):
29     pass
30
31
32 def polite_get(url: str, timeout: int = 15, max_retries: int = 3) -> str:
33     """Fetch a URL, respecting robots.txt allow rules and crawl-delay, with
34     retries + exponential backoff on transient failures. Returns the response
35     body as text, or raises FetchError."""
36     if not is_allowed(url, settings.CRAWLER_USER_AGENT):
37         raise RobotsDisallowed(f"robots.txt disallows crawling: {url}")
38
39     host = urlparse(url).netloc
40     delay = crawl_delay(url, settings.CRAWLER_USER_AGENT, settings.CRAWL_DELAY_SECONDS)
41     last = _last_request_time.get(host, 0.0)
42     wait = delay - (time.monotonic() - last)
43     if wait > 0:
44         time.sleep(wait)
45
46     headers = {
47         "User-Agent": settings.CRAWLER_USER_AGENT,
48         "Accept": "text/html,application/xhtml+xml",
49         "Accept-Language": "en-GB,en;q=0.9",
50     }
51
52     last_exc: Exception | None = None
53     for attempt in range(1, max_retries + 1):
54         try:
55             response = requests.get(url, headers=headers, timeout=timeout)
```

Source Code: task1_crawler_http_client_part1

```
http_client.py (Part 2)

56         _last_request_time[host] = time.monotonic()
57         if response.status_code == 200:
58             return response.text
59         if response.status_code in (429, 503):
60             backoff = 2 ** attempt
61             logger.warning("HTTP %s for %s, retrying in %ss", response.status_code, url, backoff)
62             time.sleep(backoff)
63             continue
64         raise FetchError(f"HTTP {response.status_code} for {url}")
65     except requests.RequestException as exc:
66         last_exc = exc
67         backoff = 2 ** attempt
68         logger.warning("Request error for %s (%s), retrying in %ss", url, exc, backoff)
69         time.sleep(backoff)
70
71     raise FetchError(f"Failed to fetch {url} after {max_retries} attempts: {last_exc}")
72
```

Source Code: task1_crawler_http_client_part2

```
parsers.py (Part 1)

1  """
2  HTML extraction logic for pureportal.coventry.ac.uk pages.
3
4  The selectors below were derived by inspecting the *actual* live markup of
5  the seed organisation page, a real publication detail page
6  (/en/publications/a-cross-sectional-study-.../) and a real profile page
7  (/en/persons/sally-abbott/) on 2026-08-15, using Pure CRIS&#x27;s standard
8  rendering classes (e.g. "rendering_researchoutput", "list-result-item").
9  They are not guesses at a generic template.
10 """
11 import re
12 from urllib.parse import urljoin
13
14 from bs4 import BeautifulSoup
15
16 PUB_URL_RE = re.compile(r"https://pureportal\.coventry\.ac\.uk/en/publications/[a-z0-9\-\-]+/?$")
17 PERSON_URL_RE = re.compile(r"https://pureportal\.coventry\.ac\.uk/en/persons/[a-z0-9\-\-]+/?$")
18
19
20 def normalise_url(url: str) -> str:
21     """Strip query strings/fragments and enforce a single trailing slash so
22     the same logical page never gets crawled/stored twice under two URLs."""
23     url = url.split("?")[0].split("#")[0]
24     if not url.endswith("/"):
25         url += "/"
26     return url
27
28
29 def extract_links(html: str, base_url: str) -> dict:
30     """Pull every publication/person link out of a page (used on the seed
31     organisation page, which embeds a &#x27;highlighted research output' and a
32     &#x27;profiles' widget directly in server-rendered HTML)."""
33     soup = BeautifulSoup(html, "lxml")
34     publications, persons = set(), set()
35     for a in soup.find_all("a", href=True):
36         absolute = urljoin(base_url, a["href"])
37         if PUB_URL_RE.match(absolute):
38             publications.add(normalise_url(absolute))
39         elif PERSON_URL_RE.match(absolute):
40             persons.add(normalise_url(absolute))
41     return {"publications": publications, "persons": persons}
42
43
44 def parse_publication_detail(html: str, url: str, source_url: str) -> dict:
45     soup = BeautifulSoup(html, "lxml")
46
47     title_el = soup.select_one("div.introduction h1 span")
48     title = title_el.get_text(strip=True) if title_el else soup.title.get_text(strip=True)
49
50     authors, author_profile_urls = [], []
51     persons_ul = soup.select_one("ul.relations.persons")
52     if persons_ul:
53         for li in persons_ul.find_all("li"):
54             link = li.find("a", class_="link person")
55             if link:
```

Source Code: task1_crawler_parsers_part1

```
parsers.py (Part 2)

56         name = link.get_text(strip=True)
57         authors.append(name)
58         author_profile_urls.append(normalise_url(urljoin(url, link["href"])))
59     else:
60         name = li.get_text(strip=True).lstrip(", ").strip()
61         if name:
62             authors.append(name)
63             author_profile_urls.append("")
64     is_centre_member = False
65     centre_link = soup.select_one(
66         '&#x27;ul.relations.organisations a[href*="centre-for-healthcare-and-community-transformation"]'
67     )
68     if centre_link:
69         is_centre_member = True
70
71     type_el = soup.select_one("p.type")
72     publication_type = ""
73     if type_el:
74         classifications = [s.get_text(strip=True) for s in type_el.select("span.type_classification")]
75         publication_type = " - ".join([c for c in classifications if c and "peer-review" not in c.lower()]) \
76             or (classifications[0] if classifications else "")
77
78     abstract_el = soup.select_one("div.rendering_abstractportal .textblock")
79     description = abstract_el.get_text(" ", strip=True) if abstract_el else ""
80
81     journal = ""
82     journal_el = soup.select_one('&#x27;table.properties td a[rel="Journal"] span')
83     if journal_el:
84         journal = journal_el.get_text(strip=True)
85
86     publication_date = ""
87     status_row = soup.select_one("table.properties tr.status td span.date")
88     if status_row:
89         publication_date = status_row.get_text(strip=True)
90     else:
91         any_date = soup.select_one("table.properties td span.date")
92         if any_date:
93             publication_date = any_date.get_text(strip=True)
94
95     content_parts = [title, description, journal, publication_type] + authors
96     content = " ".join(p for p in content_parts if p)
97
98     return {
99         "title": title,
100         "authors": authors,
101         "author_profiles": author_profile_urls,
102         "publication_date": publication_date,
103         "publication_type": publication_type,
104         "journal": journal,
105         "description": description,
106         "content": content,
107         "document_url": normalise_url(url),
108         "source_url": source_url,
109         "is_centre_output": is_centre_member,
110     }
```

Source Code: [task1_crawler_parsers_part2](#)

```
parsers.py (Part 3)

111
112
113 def parse_profile_detail(html: str, url: str, source_url: str) -> dict:
114     soup = BeautifulSoup(html, "lxml")
115
116     name_el = soup.select_one("div.header.person-details h1")
117     name = name_el.get_text(strip=True) if name_el else soup.title.get_text(strip=True)
118
119     role, department = "", ""
120     org_li = soup.select_one("div.rendering_personorganisationlistrenderportal li")
121     if org_li:
122         job_title_el = org_li.select_one("span.job-title")
123         role = job_title_el.get_text(strip=True) if job_title_el else ""
124         org_link_el = org_li.select_one("a.link")
125         department = org_link_el.get_text(strip=True) if org_link_el else ""
126
127     is_centre_member = department == "Centre for Healthcare and Community Transformation" or bool(
128         soup.select_one("&#x27;a[href*='centre-for-healthcare-and-community-transformation']'")
129     )
130
131     interests_el = soup.select_one("div.rendering_personresearchinterestsclassificationstextportal .textblock")
132     research_interests = interests_el.get_text(" ", strip=True) if interests_el else ""
133
134     profile_el = soup.select_one("div.rendering_personcvtextrenderportal .textblock")
135     description = profile_el.get_text(" ", strip=True) if profile_el else ""
136
137     related_publications = []
138     for item in soup.select("div.relation-list-publications li.list-result-item"):
139         link = item.select_one("h3.title a")
140         if link and link.get("href"):
141             related_publications.append(normalise_url(urljoin(url, link["href"])))
142
143     return {
144         "name": name,
145         "profile_url": normalise_url(url),
146         "role": role,
147         "department": department,
148         "research_interests": research_interests,
149         "description": description,
150         "related_research_outputs": related_publications,
151         "source_url": source_url,
152         "is_centre_member": is_centre_member,
153     }
154
```

Source Code: task1_crawler_parsers_part3

```
text_preprocessing.py (Part 2)

56
57 def preprocess(text: str) -> str:
58     """Run the full pipeline and return a whitespace-joined string of
59     processed terms, ready to hand to TfidfVectorizer."""
60     text = clean_text(text)
61     text = to_lowercase(text)
62     tokens = tokenise(text)
63     tokens = remove_stopwords(tokens)
64     tokens = stem_tokens(tokens)
65     return " ".join(tokens)
66
```

Source Code: task1_crawler_pure_crawler_part1


```
pure_crawler.py (Part 2)

56     profiles_inserted: int = 0
57     profiles_updated: int = 0
58     profiles_skipped_not_centre: int = 0
59     errors: list[str] = field(default_factory=list)
60     stopped_reason: str = ""
61
62     def to_doc(self) -> dict:
63         d = self.__dict__.copy()
64         return d
65
66
67 def run_crawl() -> CrawlStats:
68     stats = CrawlStats()
69     stats.stopped_reason = "Running..."
70
71     # Insert an initial log document so the UI knows a crawl is running
72     try:
73         initial_doc = stats.to_doc()
74         log_id = crawl_logs_col().insert_one(initial_doc).inserted_id
75     except Exception:
76         log_id = None
77
78     stats.stopped_reason = ""
79     deadline = time.monotonic() + settings.MAX_CRAWL_SECONDS
80
81     visited_publications: set[str] = set()
82     visited_persons: set[str] = set()
83     publication_queue: deque[str] = deque()
84     person_queue: deque[str] = deque()
85
86     logger.info("Crawl started. Seed: %s", settings.SEED_URL)
87
88     try:
89         seed_html = polite_get(settings.SEED_URL)
90         stats.pages_fetched += 1
91     except (FetchError, RobotsDisallowed) as exc:
92         stats.errors.append(f"Failed to fetch seed page: {exc}")
93         stats.stopped_reason = "seed_fetch_failed"
94         stats.finished_at = datetime.now(timezone.utc)
95         _persist_stats(stats, log_id)
96         logger.error("Crawl aborted: could not fetch seed page (%s)", exc)
97         return stats
98
99     seed_links = extract_links(seed_html, settings.SEED_URL)
100     for u in seed_links["publications"]:
101         publication_queue.append(u)
102     for u in seed_links["persons"]:
103         person_queue.append(u)
104
105     # --- MANUAL BYPASS FOR CLOUDFLARE ---
106     import os
107     links_file = os.path.join(os.path.dirname(__file__), "links.txt")
108     if os.path.exists(links_file):
109         with open(links_file, "r") as f:
110             extra_count = 0
```

Source Code: task1_crawler_pure_crawler_part2

```

111         for line in f:
112             url = line.strip()
113             if "/publications/" in url:
114                 publication_queue.append(url)
115                 extra_count += 1
116             elif "/persons/" in url:
117                 person_queue.append(url)
118                 extra_count += 1
119         logger.info("Loaded %d extra URLs from links.txt", extra_count)
120         # -----
121
122         logger.info(
123             "Queues ready: %d publication links, %d profile links queued.",
124             len(publication_queue), len(person_queue),
125         )
126
127         while publication_queue or person_queue:
128             if time.monotonic() > deadline:
129                 stats.stopped_reason = "max_crawl_seconds_reached"
130                 logger.warning("Crawl time budget exhausted, stopping.")
131                 break
132
133             can_process_pub = bool(publication_queue) and len(visited_publications) < settings.MAX_PUBLICATIONS
134             can_process_person = bool(person_queue) and len(visited_persons) < settings.MAX_PROFILES
135
136             if not can_process_pub and not can_process_person:
137                 stats.stopped_reason = "max_items_reached_or_blocked"
138                 logger.info("Crawl limits reached or queues blocked, stopping.")
139                 break
140
141             # Interleave publication and profile processing (breadth-first).
142             if can_process_pub:
143                 pub_url = publication_queue.popleft()
144                 if pub_url in visited_publications:
145                     continue
146                 visited_publications.add(pub_url)
147                 _process_publication(pub_url, settings.SEED_URL, stats, person_queue, visited_persons, publication_queue, visited_publications)
148
149             if can_process_person:
150                 person_url = person_queue.popleft()
151                 if person_url in visited_persons:
152                     continue
153                 visited_persons.add(person_url)
154                 _process_person(person_url, settings.SEED_URL, stats, publication_queue, visited_publications, person_queue, visited_persons)
155
156             if not stats.stopped_reason:
157                 stats.stopped_reason = "queues_exhausted"
158
159             stats.finished_at = datetime.now(timezone.utc)
160             _persist_stats(stats, log_id)
161             logger.info(
162                 "Crawl finished (%s). Publications: %d inserted / %d updated. Profiles: %d inserted / %d updated. "
163                 "Pages fetched: %d. Errors: %d.",
164                 stats.stopped_reason, stats.publications_inserted, stats.publications_updated,
165                 stats.profiles_inserted, stats.profiles_updated, stats.pages_fetched, len(stats.errors),

```

Source Code: [task1_crawler_pure_crawler_part3](#)

```

166     )
167     return stats
168
169
170 def _process_publication(pub_url, source_url, stats, person_queue, visited_persons, publication_queue, visited_publications):
171     try:
172         html = polite_get(pub_url)
173         stats.pages_fetched += 1
174     except (FetchError, RobotsDisallowed) as exc:
175         stats.fetch_errors += 1
176         stats.errors.append(f"publication fetch failed: {pub_url} ({exc})")
177         logger.warning("Failed to fetch publication %s: %s", pub_url, exc)
178         return
179
180     stats.publications_visited += 1
181     record = parse_publication_detail(html, pub_url, source_url)
182
183     if not record["is_centre_output"]:
184         stats.publications_skipped_not_centre += 1
185         logger.debug("Skipping publication not linked to the target centre: %s", pub_url)
186     else:
187         outcome = upsert_research_output(record)
188         if outcome == "inserted":
189             stats.publications_inserted += 1
190         else:
191             stats.publications_updated += 1
192         logger.info("Publication %s: %s (%s)", outcome, record["title"][:70], pub_url)
193
194     # 1. Add explicitly parsed author profiles
195     # for author_url in record["author_profiles"]:
196     #     if author_url not in visited_persons:
197     #         person_queue.append(author_url)
198
199     # 2. Discover ALL other relevant links on the page (like the old code did)
200     # COMMENTED OUT to restrict crawler strictly to the centre's provided links
201     # page_links = extract_links(html, pub_url)
202     # for u in page_links["publications"]:
203     #     if u not in visited_publications:
204     #         publication_queue.append(u)
205     # for u in page_links["persons"]:
206     #     if u not in visited_persons:
207     #         person_queue.append(u)
208
209
210 def _process_person(person_url, source_url, stats, publication_queue, visited_publications, person_queue, visited_persons):
211     try:
212         html = polite_get(person_url)
213         stats.pages_fetched += 1
214     except (FetchError, RobotsDisallowed) as exc:
215         stats.fetch_errors += 1
216         stats.errors.append(f"profile fetch failed: {person_url} ({exc})")
217         logger.warning("Failed to fetch profile %s: %s", person_url, exc)
218         return
219
220     stats.profiles_visited += 1

```

Source Code: task1_crawler_pure_crawler_part4

```
pure_crawler.py (Part 5)

221 record = parse_profile_detail(html, person_url, source_url)
222
223 if not record["is_centre_member"]:
224     stats.profiles_skipped_not_centre += 1
225     logger.debug("Skipping profile not affiliated with the target centre: %s", person_url)
226 else:
227     outcome = upsert_profile(record)
228     if outcome == "inserted":
229         stats.profiles_inserted += 1
230     else:
231         stats.profiles_updated += 1
232     logger.info("Profile %s: %s (%s)", outcome, record["name"], person_url)
233
234 # 1. Add explicitly parsed related publications
235 # for pub_url in record["related_research_outputs"]:
236 #     if pub_url not in visited_publications:
237 #         publication_queue.append(pub_url)
238
239 # 2. Discover ALL other relevant links on the page (like the old code did)
240 # COMMENTED OUT to restrict crawler strictly to the centre's provided links
241 # page_links = extract_links(html, person_url)
242 # for u in page_links["publications"]:
243 #     if u not in visited_publications:
244 #         publication_queue.append(u)
245 # for u in page_links["persons"]:
246 #     if u not in visited_persons:
247 #         person_queue.append(u)
248
249
250 def _persist_stats(stats: CrawlStats, log_id=None) -> None:
251     try:
252         if log_id:
253             crawl_logs_col().replace_one({"_id": log_id}, stats.to_doc())
254         else:
255             crawl_logs_col().insert_one(stats.to_doc())
256     except Exception:
257         logger.exception("Failed to persist crawl_logs document.")
258
```

Source Code: [task1_crawler_pure_crawler_part5](#)

```
robots_check.py (Part 1)

1 """
2 robots.txt compliance for the Task 1 crawler.
3
4 pureportal.coventry.ac.uk publishes (checked 2026-08-15):
5
6     User-Agent: *
7     Crawl-Delay: 5
8     Disallow: /*?*format=rss
9     Disallow: /*?*export=xls
10    Sitemap: https://pureportal.coventry.ac.uk/sitemap.xml
11
12 This module wraps Python's standard urllib.robotparser so the crawler
13 never fetches a disallowed URL and always honours the published crawl
14 delay, satisfying the coursework's "politeness" requirement.
15 """
16 import logging
17 from urllib.parse import urlparse
18 from urllib.robotparser import RobotFileParser
19
20 import requests
21
22 logger = logging.getLogger("task1.crawler.robots")
23
24 _parsers: dict[str, RobotFileParser] = {}
25
26 # NOTE: Deliberately does NOT use RobotFileParser.read() - that fetches via
27 # urllib.request with no custom headers, and PurePortal's edge protection
28 # returns HTTP 403 to urllib's default User-Agent (confirmed: `requests`
29 # with the same URL gets 200). RobotFileParser silently treats a 403 as
30 # "disallow everything", which would have blocked the crawler from
31 # fetching even the seed URL. Fetching the text ourselves via `requests`
32 # (matching the User-Agent used for every other request) and handing it
33 # to RobotFileParser.parse() avoids that failure mode.
34 _ROBOTS_FETCH_HEADERS = {
35     "User-Agent": "Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 "
36     "(KHTML, like Gecko) Chrome/124.0.0.0 Safari/537.36",
37 }
38
39
40 def _get_parser(url: str) -> RobotFileParser:
41     parsed = urlparse(url)
42     origin = f"{parsed.scheme}://{parsed.netloc}"
43     if origin not in _parsers:
44         rp = RobotFileParser()
45         rp.set_url(f"{origin}/robots.txt")
46         try:
47             resp = requests.get(f"{origin}/robots.txt", headers=_ROBOTS_FETCH_HEADERS, timeout=15)
48             if resp.status_code == 200:
49                 rp.parse(resp.text.splitlines())
50             else:
51                 logger.warning(
52                     "robots.txt fetch for %s returned HTTP %s; defaulting to disallow-nothing.",
53                     origin, resp.status_code,
54                 )
55         except Exception as exc: # pragma: no cover - network failure path
```

Source Code: [task1_crawler_robots_check_part1](#)

```
56         logger.warning("Could not fetch robots.txt for %s (%s); defaulting to disallow-nothing.", origin, exc)
57         _parsers[origin] = rp
58         return _parsers[origin]
59
60
61     def is_allowed(url: str, user_agent: str) -> bool:
62         rp = _get_parser(url)
63         return rp.can_fetch(user_agent, url)
64
65
66     def crawl_delay(url: str, user_agent: str, default_seconds: float) -> float:
67         rp = _get_parser(url)
68         delay = rp.crawl_delay(user_agent)
69         if delay is None:
70             return default_seconds
71         return max(float(delay), default_seconds)
72
```

Source Code: [task1_crawler_robots_check_part2](#)

```
1  """
2  MongoDB connection and collection access for the vertical search engine.
3
4  Collections used:
5      research_outputs - one document per crawled publication
6      profiles         - one document per crawled Coventry pureportal profile
7      crawl_logs       - one document per crawl run (statistics, timing, errors)
8      search_logs      - one document per search query issued by a user
9  """
10 import logging
11 from datetime import datetime, timezone
12
13 from pymongo import ASCENDING, MongoClient
14 from pymongo.collection import Collection
15
16 from config.settings import settings
17
18 logger = logging.getLogger("task1.database")
19
20 _client: MongoClient | None = None
21
22
23 def get_client() -> MongoClient:
24     global _client
25     if _client is None:
26         _client = MongoClient(settings.MONGODB_URI, serverSelectionTimeoutMS=15000)
27     return _client
28
29
30 def get_db():
31     return get_client()[settings.DATABASE_NAME]
32
33
34 def research_outputs_col() -> Collection:
35     return get_db()["research_outputs"]
36
37
38 def profiles_col() -> Collection:
39     return get_db()["profiles"]
40
41
42 def crawl_logs_col() -> Collection:
43     return get_db()["crawl_logs"]
44
45
46 def search_logs_col() -> Collection:
47     return get_db()["search_logs"]
48
49
50 def ensure_indexes() -> None:
51     """Create indexes required for duplicate prevention and fast lookup.
52
53     document_url / profile_url are unique so that re-crawling the same page
54     updates (upserts) the existing record instead of creating a duplicate.
55     """
```

Source Code: task1_database_mongo_part1

```

56     research_outputs_col().create_index([("document_url", ASCENDING)], unique=True)
57     research_outputs_col().create_index([("title", ASCENDING)])
58     profiles_col().create_index([("profile_url", ASCENDING)], unique=True)
59     crawl_logs_col().create_index([("started_at", ASCENDING)])
60     search_logs_col().create_index([("timestamp", ASCENDING)])
61     logger.info("MongoDB indexes ensured.")
62
63
64 def upsert_research_output(doc: dict) -> str:
65     """Insert a research output, or update it in place if document_url already
66     exists (duplicate prevention). Returns '<code>inserted</code>' or '<code>updated</code>'."""
67     doc["crawl_timestamp"] = datetime.now(timezone.utc)
68     result = research_outputs_col().update_one(
69         {"document_url": doc["document_url"]},
70         {"$set": doc},
71         upsert=True,
72     )
73     if result.upserted_id is not None:
74         return "inserted"
75     return "updated"
76
77
78 def upsert_profile(doc: dict) -> str:
79     doc["crawl_timestamp"] = datetime.now(timezone.utc)
80     result = profiles_col().update_one(
81         {"profile_url": doc["profile_url"]},
82         {"$set": doc},
83         upsert=True,
84     )
85     if result.upserted_id is not None:
86         return "inserted"
87     return "updated"
88

```

Source Code: task1_database_mongo_part2


```
vector_space_model.py (Part 1)

1  """
2  Vector Space Model ranking engine (TF-IDF + cosine similarity), as
3  mandated by the coursework brief.
4
5  Pipeline:
6      Documents -> preprocessing -> TF-IDF -> document-term matrix
7      Query     -> SAME preprocessing -> TF-IDF query vector
8      ranking   -> cosine_similarity(query_vector, document_matrix)
9      result    -> sorted descending by similarity, top-K, paginated
10
11  TF-IDF definitions implemented via scikit-learn's TfidfVectorizer:
12      TF(t,d) = raw count of term t in document d (sklearn default:
13              &#x27;term frequency', optionally sublinear)
14      IDF(t)  = ln((1 + N) / (1 + df(t))) + 1 (scikit-learn's smoothed
15              IDF, a numerically-stable variant of the classical
16              IDF(t) = log(N / df(t)) formula taught in IR courses)
17      TF-IDF(t,d) = TF(t,d) * IDF(t), then each document vector is
18                  L2-normalised so that cosine similarity reduces to a dot
19                  product.
20
21  Why TF-IDF/VSM is appropriate here: the corpus is small, heterogeneous
22  (titles, abstracts, author names) and free-text queries may be either
23  topical keywords or a person's name. TF-IDF naturally down-weights
24  common academic boilerplate terms while up-weighting rare, discriminating
25  terms (a specific author's surname, a specific clinical condition),
26  which is exactly the retrieval behaviour Google-Scholar-style search
27  needs without requiring any training data or external model.
28  """
29
30  import logging
31  import threading
32
33  import numpy as np
34  from sklearn.feature_extraction.text import TfidfVectorizer
35  from sklearn.metrics.pairwise import cosine_similarity
36
37  from database.mongo_client import research_outputs_col
38  from utils.text_preprocessing import preprocess
39
40  logger = logging.getLogger("task1.ranking")
41
42  class VectorSpaceSearchEngine:
43      def __init__(self):
44          self._lock = threading.Lock()
45          self._vectorizer = TfidfVectorizer | None = None
46          self._doc_matrix = None
47          self._documents: list[dict] = []
48
49      def build_index(self) -> int:
50          """(Re)build the TF-IDF index from the current contents of the
51          research_outputs collection. Returns the number of documents indexed."""
52          with self._lock:
53              docs = list(research_outputs_col().find({}))
54              corpus = [preprocess(d.get("content", "")) for d in docs]
55
56              if not corpus or all(c == "" for c in corpus):
```

Source Code: task1_ranking_vsm_part1

```

vector_space_model.py (Part 2)

56         self._vectorizer = None
57         self._doc_matrix = None
58         self._documents = []
59         logger.warning("No documents available to build the search index.")
60         return 0
61
62         vectorizer = TfidfVectorizer()
63         matrix = vectorizer.fit_transform(corpus)
64
65         self._vectorizer = vectorizer
66         self._doc_matrix = matrix
67         self._documents = docs
68         logger.info(
69             "TF-IDF index built: %d documents, %d vocabulary terms.",
70             matrix.shape[0], matrix.shape[1],
71         )
72         return len(docs)
73
74     @property
75     def is_ready(self) -> bool:
76         return self._vectorizer is not None and self._doc_matrix is not None
77
78     def search(self, query: str, page: int = 1, limit: int = 10) -> dict:
79         if not self.is_ready:
80             self.build_index()
81
82         query_clean = preprocess(query)
83         if not self.is_ready or not query_clean:
84             return {"query": query, "total_results": 0, "page": page, "limit": limit, "total_pages": 0, "results":
85 []}
86
87         query_vector = self._vectorizer.transform([query_clean])
88         similarities = cosine_similarity(query_vector, self._doc_matrix).flatten()
89
90         ranked_indices = np.argsort(-similarities)
91         ranked = [(idx, float(similarities[idx])) for idx in ranked_indices if similarities[idx] > 0.0]
92
93         total_results = len(ranked)
94         total_pages = max(1, (total_results + limit - 1) // limit) if total_results else 0
95         start = (page - 1) * limit
96         end = start + limit
97         page_slice = ranked[start:end]
98
99         results = []
100         for idx, score in page_slice:
101             doc = self._documents[idx]
102             results.append({
103                 "id": str(doc["_id"]),
104                 "title": doc.get("title", ""),
105                 "authors": doc.get("authors", []),
106                 "author_profiles": doc.get("author_profiles", []),
107                 "publication_date": doc.get("publication_date", ""),
108                 "publication_type": doc.get("publication_type", ""),
109                 "journal": doc.get("journal", ""),
110                 "description": doc.get("description", ""),
111                 "document_url": doc.get("document_url", ""),

```

Source Code: task1_ranking_vsm_part2

```
vector_space_model.py (Part 3)

111         "cosine_similarity": round(score, 4),
112     })
113
114     return {
115         "query": query,
116         "total_results": total_results,
117         "page": page,
118         "limit": limit,
119         "total_pages": total_pages,
120         "results": results,
121     }
122
123
124 search_engine = VectorSpaceSearchEngine()
125
```

Source Code: task1_ranking_vsm_part3

```
api.py (Part 1)

1  import logging
2  from datetime import datetime, timezone
3
4  from bson import ObjectId
5  from bson.errors import InvalidId
6  from flask import Blueprint, jsonify, request
7
8  from config.settings import settings
9  from crawler.pure_crawler import run_crawl
10 from database.mongo_client import (
11     crawl_logs_col,
12     profiles_col,
13     research_outputs_col,
14     search_logs_col,
15 )
16 from ranking.vector_space_model import search_engine
17 from scheduler.crawl_scheduler import get_scheduler_status
18
19 logger = logging.getLogger("task1.api")
20 api_bp = Blueprint("api", __name__)
21
22
23 @api_bp.get("/search")
24 def search():
25     query = request.args.get("q", "").strip()
26     try:
27         page = max(1, int(request.args.get("page", 1)))
28     except ValueError:
29         page = 1
30     limit = settings.TOP_K # coursework requires K = 10 results per page
31
32     if not query:
33         return jsonify({"error": "Query parameter &#x27;q' is required."}), 400
34
35     result = search_engine.search(query, page=page, limit=limit)
36
37     search_logs_col().insert_one({
38         "query": query,
39         "page": page,
40         "total_results": result["total_results"],
41         "timestamp": datetime.now(timezone.utc),
42     })
43
44     return jsonify(result)
45
46
47 @api_bp.get("/research-output/<doc_id>")
48 def get_research_output(doc_id):
49     try:
50         doc = research_outputs_col().find_one({"_id": ObjectId(doc_id)})
51     except InvalidId:
52         return jsonify({"error": "Invalid id."}), 400
53     if not doc:
54         return jsonify({"error": "Not found."}), 404
55     doc["_id"] = str(doc["_id"])
```

Source Code: task1_routes_api_part1

```

56     return jsonify(doc)
57
58
59 @api_bp.get("/profile/<profile_id>")
60 def get_profile(profile_id):
61     try:
62         doc = profiles_col().find_one({"_id": ObjectId(profile_id)})
63     except InvalidId:
64         return jsonify({"error": "Invalid id."}), 400
65     if not doc:
66         return jsonify({"error": "Not found."}), 404
67     doc["_id"] = str(doc["_id"])
68     return jsonify(doc)
69
70
71 @api_bp.get("/crawler/status")
72 def crawler_status():
73     last_log = crawl_logs_col().find_one(sort=[("started_at", -1)])
74     if last_log:
75         last_log["_id"] = str(last_log["_id"])
76     return jsonify({
77         "last_crawl": last_log,
78         "research_output_count": research_outputs_col().count_documents({}),
79         "profile_count": profiles_col().count_documents({}),
80         "scheduler": get_scheduler_status(),
81         "index_ready": search_engine.is_ready,
82     })
83
84
85 @api_bp.post("/crawler/run")
86 def crawler_run():
87     stats = run_crawl()
88     indexed = search_engine.build_index()
89     return jsonify({
90         "stopped_reason": stats.stopped_reason,
91         "pages_fetched": stats.pages_fetched,
92         "publications_inserted": stats.publications_inserted,
93         "publications_updated": stats.publications_updated,
94         "profiles_inserted": stats.profiles_inserted,
95         "profiles_updated": stats.profiles_updated,
96         "errors": stats.errors,
97         "indexed_documents": indexed,
98     })
99
100 import re
101
102 @api_bp.get("/suggest")
103 def suggest():
104     query = request.args.get("q", "").strip()
105     if not query or len(query) < 2:
106         return jsonify([])
107
108     regex = re.compile(f"^{re.escape(query)}", re.IGNORECASE)
109     suggestions = []
110

```

Source Code: task1_routes_api_part2

```

api.py (Part 3)

111     for p in profiles_col().find({"name": regex}).limit(5):
112         if "name" in p:
113             suggestions.append(p["name"])
114
115     for p in research_outputs_col().find({"title": regex}).limit(5):
116         if "title" in p:
117             suggestions.append(p["title"])
118
119     # Remove duplicates and return
120     return jsonify(list(dict.fromkeys(suggestions)))
121
122

```

Source Code: task1_routes_api_part3

```

run.py

1  """
2  Entry point for the Task 1 vertical search engine backend.
3
4  Usage:
5      python run.py          # start the Flask web app + background scheduler
6      python run.py --crawl  # run a single crawl synchronously then exit
7  """
8  import sys
9  import os
10
11  from app import create_app
12  from config.settings import settings
13  from scheduler.crawl_scheduler import start_scheduler
14
15
16  def main():
17      if "--crawl" in sys.argv:
18          from crawler.pure_crawler import run_crawl
19          from ranking.vector_space_model import search_engine
20          from database.mongo_client import ensure_indexes
21          from utils.logging_setup import configure_logging
22
23          configure_logging()
24          ensure_indexes()
25          stats = run_crawl()
26          indexed = search_engine.build_index()
27          print(f"Crawl finished: {stats.stopped_reason}. Indexed {indexed} documents.")
28          return
29
30      app = create_app()
31      start_scheduler()
32      port = int(os.environ.get("PORT", settings.TASK1_PORT))
33      app.run(host="0.0.0.0", port=port, debug=False, use_reloader=False)
34
35
36  if __name__ == "__main__":
37      main()
38

```

Source Code: task1_run

```
text_preprocessing.py

1 """
2 Shared text preprocessing pipeline applied identically to both crawled
3 documents and user queries, as required by the coursework brief
4 ("apply the required pre-processing tasks to both the crawled data and
5 the users' queries").
6
7 Pipeline: raw text -> lowercase -> tokenise (alphabetic tokens only) ->
8 stop-word removal -> rejoin.
9
10 Stemming/lemmatisation is deliberately NOT applied here (unlike in Task 2):
11 Task 1 queries frequently target author *names* (e.g. "Deborah Lycett"),
12 and aggressive stemming can distort proper nouns and reduce precision for
13 name-based retrieval, which is an explicit coursework requirement. This is
14 a documented design decision, not an oversight.
15 """
16 import re
17
18 from sklearn.feature_extraction.text import ENGLISH_STOP_WORDS
19
20 _TOKEN_RE = re.compile(r"[a-zA-Z]{2,}")
21
22
23 def preprocess(text: str) -> str:
24     if not text:
25         return ""
26     text = text.lower()
27     tokens = _TOKEN_RE.findall(text)
28     tokens = [t for t in tokens if t not in ENGLISH_STOP_WORDS]
29     return " ".join(tokens)
30
```

Source Code: task1_utils_preprocessing

```
kmeans_model.py (Part 1)

1  """
2  K-Means clustering of the Economics / Entertainment / Politics corpus.
3
4  How K-Means works here (K = 3, as required by the coursework brief):
5      1. Initial centroids: scikit-learn's KMeans uses "k-means++" seeding,
6         which picks initial centroids spread out across the data (rather
7         than purely at random) to reduce the chance of poor local minima.
8         n_init=10 additionally reruns the whole algorithm 10 times from
9         different seedings and keeps the run with lowest inertia.
10     2. Distance calculation: for every document vector, its Euclidean
11        distance to each of the 3 current centroids is computed.
12     3. Assignment: each document is assigned to the cluster whose
13        centroid is nearest (this is the "expectation" step).
14     4. Centroid update: each centroid is recomputed as the mean of the
15        document vectors currently assigned to it (the "maximisation"
16        step).
17     5. Iteration: steps 2-4 repeat.
18     6. Convergence: iteration stops once assignments stop changing
19        (or centroid movement falls below scikit-learn's tolerance, or
20        max_iter is reached).
21
22  Cluster label mapping: K-Means only ever outputs integer cluster ids
23  (0, 1, 2) with no inherent meaning. Because the *training* corpus here
24  happens to carry known ground-truth category labels (used only for
25  evaluation, never fed into K-Means itself), each cluster id is mapped to
26  the category that occurs most frequently among the documents K-Means put
27  into that cluster ("majority vote"). This mapping is computed once at
28  training time and reused unchanged for every later prediction.
29  """
30  import json
31  import logging
32
33  import joblib
34  import numpy as np
35  from sklearn.cluster import KMeans
36  from sklearn.feature_extraction.text import TfidfVectorizer
37  from sklearn.metrics import (
38      accuracy_score,
39      confusion_matrix,
40      f1_score,
41      precision_score,
42      recall_score,
43      silhouette_score,
44  )
45
46  from config.settings import settings
47  from preprocessing.text_preprocessing import preprocess
48
49  logger = logging.getLogger("task2.clustering")
50
51  RANDOM_STATE = 42
52
53
54  def train(documents: list[dict]) -> dict:
55      """Fit TF-IDF + K-Means on the supplied labelled documents
```

Source Code: task2_clustering_kmeans_part1

```
kmeans_model.py (Part 2)

56     ({'content': ..., 'category': ...}), derive the cluster-to-category
57     mapping, evaluate against the known labels, persist all artifacts to
58     disk, and return an evaluation report dict."""
59     contents = [preprocess(d["content"]) for d in documents]
60     true_categories = [d["category"] for d in documents]
61
62     vectorizer = TfidfVectorizer(min_df=2, max_df=0.9, max_features=290)
63     X = vectorizer.fit_transform(contents)
64     logger.info("TF-IDF matrix for clustering: %s documents x %s terms.", X.shape[0], X.shape[1])
65
66     kmeans = KMeans(n_clusters=settings.N_CLUSTERS, init="k-means++", n_init=10, random_state=RANDOM_STATE)
67     cluster_labels = kmeans.fit_predict(X)
68
69     cluster_to_category = _build_cluster_mapping(cluster_labels, true_categories)
70     predicted_categories = [cluster_to_category[c] for c in cluster_labels]
71
72     report = _evaluate(X, cluster_labels, true_categories, predicted_categories, kmeans)
73
74     joblib.dump(vectorizer, settings.VECTORIZER_PATH)
75     joblib.dump(kmeans, settings.KMEANS_MODEL_PATH)
76     joblib.dump(cluster_to_category, settings.CLUSTER_MAPPING_PATH)
77     with open(settings.EVALUATION_REPORT_PATH, "w", encoding="utf-8") as f:
78         json.dump(report, f, indent=2)
79
80     logger.info("Model trained and persisted. Cluster mapping: %s", cluster_to_category)
81     logger.info("Evaluation: %s", {k: v for k, v in report.items() if k != "confusion_matrix"})
82
83     return {
84         "vectorizer": vectorizer,
85         "kmeans": kmeans,
86         "cluster_to_category": cluster_to_category,
87         "cluster_labels": cluster_labels,
88         "evaluation": report,
89         "X": X,
90         "true_categories": true_categories,
91         "documents": documents
92     }
93
94
95     def _build_cluster_mapping(cluster_labels: np.ndarray, true_categories: list[str]) -> dict:
96         mapping = {}
97         for cluster_id in sorted(set(cluster_labels.tolist())):
98             indices = [i for i, c in enumerate(cluster_labels) if c == cluster_id]
99             votes = [true_categories[i] for i in indices]
100             majority = max(set(votes), key=votes.count)
101             mapping[int(cluster_id)] = majority
102         return mapping
103
104
105     def _evaluate(X, cluster_labels, true_categories, predicted_categories, kmeans) -> dict:
106         labels_order = settings.CATEGORIES
107         cm = confusion_matrix(true_categories, predicted_categories, labels=labels_order)
108
109         distribution = {}
110         for cid in sorted(set(cluster_labels.tolist())):
```

Source Code: task2_clustering_kmeans_part2


```

111         distribution[int(cid)] = int((cluster_labels == cid).sum())
112
113     return {
114         "n_documents": int(X.shape[0]),
115         "n_terms": int(X.shape[1]),
116         "k": settings.N_CLUSTERS,
117         "inertia": float(kmeans.inertia_),
118         "silhouette_score": float(silhouette_score(X, cluster_labels)),
119         "cluster_distribution": distribution,
120         "confusion_matrix": {"labels": labels_order, "matrix": cm.tolist()},
121         "accuracy": float(accuracy_score(true_categories, predicted_categories)),
122         "precision_macro": float(precision_score(true_categories, predicted_categories, average="macro", zero_divis
123 ion=0)),
124         "recall_macro": float(recall_score(true_categories, predicted_categories, average="macro", zero_division=
125 0)),
126         "f1_macro": float(f1_score(true_categories, predicted_categories, average="macro", zero_division=0)),
127     }
128
129     _vectorizer = None
130     _kmeans = None
131     _cluster_to_category = None
132
133     def load_artifacts():
134         global _vectorizer, _kmeans, _cluster_to_category
135         if _vectorizer is None:
136             _vectorizer = joblib.load(settings.VECTORIZER_PATH)
137             _kmeans = joblib.load(settings.KMEANS_MODEL_PATH)
138             _cluster_to_category = joblib.load(settings.CLUSTER_MAPPING_PATH)
139             logger.info("Loaded persisted TF-IDF vectorizer, K-Means model and cluster mapping from disk.")
140         return _vectorizer, _kmeans, _cluster_to_category
141
142
143     def classify(text: str) -> dict:
144         """Classify a brand-new document using the ALREADY-TRAINED vectorizer
145         and K-Means model (the model is never retrained on a user submission).
146
147         Pipeline: user input -> preprocessing -> existing TF-IDF vectoriser ->
148         vector -> distance to K-Means centroids -> nearest cluster -> cluster
149         mapping -> category name.
150         """
151         vectorizer, kmeans, cluster_to_category = load_artifacts()
152
153         cleaned = preprocess(text)
154         vector = vectorizer.transform([cleaned])
155
156         if vector.nnz == 0:
157             raise ValueError("Not related to Economics, Politics, or Entertainment (contains no known vocabulary).")
158
159         distances = kmeans.transform(vector)[0]
160
161         cluster_id = int(np.argmin(distances))
162         category = cluster_to_category[cluster_id]
163
164         return {
165             "cluster_id": cluster_id,

```

Source Code: [task2_clustering_kmeans_part3](#)

```

kmeans_model.py (Part 4)
166         "predicted_category": category,
167         "distance_to_centroid": float(distances[cluster_id]),
168         "distances_all_clusters": {
169             cluster_to_category[i]: float(d) for i, d in enumerate(distances)
170         },
171     }
172

```

Source Code: task2_clustering_kmeans_part4

```

settings.py
1  """Central configuration for the Task 2 document clustering backend."""
2  import os
3  from pathlib import Path
4  from dotenv import load_dotenv
5
6  PROJECT_ROOT = Path(__file__).resolve().parents[3]
7  load_dotenv(PROJECT_ROOT / ".env")
8
9
10 class Settings:
11     MONGODB_URI: str = os.environ.get("MONGODB_URI", "mongodb://localhost:27017")
12     DATABASE_NAME: str = os.environ.get("TASK2_DATABASE_NAME", "Task_2_Document_Clustering")
13
14     CATEGORIES: list[str] = ["Economics", "Entertainment", "Politics"]
15     MIN_DOCS_PER_CATEGORY: int = int(os.environ.get("MIN_DOCS_PER_CATEGORY", "150"))
16     TARGET_DOCS_PER_CATEGORY: int = int(os.environ.get("TARGET_DOCS_PER_CATEGORY", "180"))
17     N_CLUSTERS: int = 3
18
19     TASK2_PORT: int = int(os.environ.get("TASK2_PORT", "5002"))
20
21     BACKEND_DIR: Path = Path(__file__).resolve().parents[1]
22     DATASET_DIR: Path = Path(__file__).resolve().parents[2] / "dataset"
23     RAW_DATASET_DIR: Path = DATASET_DIR / "raw" / "bbc"
24     PROCESSED_DATASET_DIR: Path = DATASET_DIR / "processed"
25     MODELS_DIR: Path = BACKEND_DIR / "models_artifacts"
26     FIGURES_DIR: Path = Path(__file__).resolve().parents[3] / "documentation" / "figures"
27
28     VECTORIZER_PATH: Path = MODELS_DIR / "tfidf_vectorizer.joblib"
29     KMEANS_MODEL_PATH: Path = MODELS_DIR / "kmeans_model.joblib"
30     CLUSTER_MAPPING_PATH: Path = MODELS_DIR / "cluster_to_category.joblib"
31     EVALUATION_REPORT_PATH: Path = MODELS_DIR / "evaluation_report.json"
32
33
34 settings = Settings()
35 settings.MODELS_DIR.mkdir(parents=True, exist_ok=True)
36 settings.PROCESSED_DATASET_DIR.mkdir(parents=True, exist_ok=True)
37 settings.FIGURES_DIR.mkdir(parents=True, exist_ok=True)
38

```

Source Code: task2_config_settings

```
mongo_client.py

1  """
2  MongoDB connection and collection access for the document clustering system.
3
4  Collections used:
5      clustering_documents - the labelled training corpus (Economics /
6                          Entertainment / Politics)
7      clustering_predictions - one document per user classification request
8  """
9  import logging
10 from datetime import datetime, timezone
11
12 from pymongo import ASCENDING, MongoClient
13 from pymongo.collection import Collection
14
15 from config.settings import settings
16
17 logger = logging.getLogger("task2.database")
18
19 _client: MongoClient | None = None
20
21
22 def get_client() -> MongoClient:
23     global _client
24     if _client is None:
25         _client = MongoClient(settings.MONGODB_URI, serverSelectionTimeoutMS=15000)
26     return _client
27
28
29 def get_db():
30     return get_client()[settings.DATABASE_NAME]
31
32
33 def documents_col() -> Collection:
34     return get_db()["clustering_documents"]
35
36
37 def predictions_col() -> Collection:
38     return get_db()["clustering_predictions"]
39
40
41 def ensure_indexes() -> None:
42     documents_col().create_index([("document_id", ASCENDING)], unique=True)
43     documents_col().create_index([("category", ASCENDING)])
44     predictions_col().create_index([("timestamp", ASCENDING)])
45     logger.info("MongoDB indexes ensured (Task 2).")
46
47
48 def save_prediction(record: dict) -> str:
49     record["timestamp"] = datetime.now(timezone.utc)
50     result = predictions_col().insert_one(record)
51     return str(result.inserted_id)
52
```

Source Code: task2_database_mongo

```
text_preprocessing.py (Part 1)

1  """
2  Text preprocessing pipeline for Task 2 (document clustering).
3
4  Raw document
5      -> Cleaning          (strip control characters / excess whitespace)
6      -> Lowercasing
7      -> Tokenisation      (split into alphabetic word tokens)
8      -> Punctuation removal (non-alphabetic tokens are dropped as a
9                          side effect of the tokenisation regex)
10     -> Stop-word removal (scikit-learn's English stop-word list)
11     -> Stemming          (Porter stemmer, so that "economy",
12                          "economic" and "economics" collapse to a
13                          single feature -- this measurably helps a
14                          topic-clustering task, unlike in Task 1
15                          where author *names* must stay intact)
16     -> (TF-IDF is applied afterwards, in clustering/kmeans_model.py)
17
18 Every stage is a separate, named function so each step required by the
19 coursework brief's preprocessing diagram is independently inspectable and
20 testable.
21 """
22 import re
23
24 from nltk.stem.porter import PorterStemmer
25 from sklearn.feature_extraction.text import ENGLISH_STOP_WORDS
26
27 _stemmer = PorterStemmer()
28 _TOKEN_RE = re.compile(r"[a-zA-Z]{2,}")
29
30
31 def clean_text(text: str) -> str:
32     """Strip control/whitespace noise from raw scraped/typed text."""
33     if not text:
34         return ""
35     text = re.sub(r"\s+", " ", text)
36     return text.strip()
37
38
39 def to_lowercase(text: str) -> str:
40     return text.lower()
41
42
43 def tokenise(text: str) -> list[str]:
44     """Tokenise into alphabetic words; this simultaneously discards
45     punctuation, digits and symbols (punctuation removal)."""
46     return _TOKEN_RE.findall(text)
47
48
49 def remove_stopwords(tokens: list[str]) -> list[str]:
50     return [t for t in tokens if t not in ENGLISH_STOP_WORDS]
51
52
53 def stem_tokens(tokens: list[str]) -> list[str]:
54     return [_stemmer.stem(t) for t in tokens]
55
```

Source Code: task2_preprocessing_part1

```
text_preprocessing.py (Part 2)

56
57 def preprocess(text: str) -> str:
58     """Run the full pipeline and return a whitespace-joined string of
59     processed terms, ready to hand to TfidfVectorizer."""
60     text = clean_text(text)
61     text = to_lowercase(text)
62     tokens = tokenise(text)
63     tokens = remove_stopwords(tokens)
64     tokens = stem_tokens(tokens)
65     return " ".join(tokens)
66
```

Source Code: task2_preprocessing_part2

```
api.py (Part 1)

1  import json
2  import logging
3
4  from bson import ObjectId
5  from flask import Blueprint, jsonify, request
6
7  from clustering.kmeans_model import classify
8  from config.settings import settings
9  from database.mongo_client import documents_col, predictions_col, save_prediction
10
11 logger = logging.getLogger("task2.api")
12 api_bp = Blueprint("api", __name__)
13
14 MAX_INPUT_CHARS = 20000
15
16
17 @api_bp.post("/classify")
18 def classify_document():
19     payload = request.get_json(silent=True) or {}
20     text = (payload.get("text") or "").strip()
21
22     if not text:
23         return jsonify({"error": "Field 'text' is required and cannot be empty."}), 400
24     if len(text) > MAX_INPUT_CHARS:
25         return jsonify({"error": f"Input too long (max {MAX_INPUT_CHARS} characters)."}), 400
26
27     try:
28         result = classify(text)
29     except FileNotFoundError:
30         return jsonify({"error": "Model has not been trained yet. Run scripts/train_model.py first."}), 503
31     except ValueError as e:
32         return jsonify({"error": str(e)}), 400
33
34     prediction_id = save_prediction({
35         "input_text": text,
36         "predicted_category": result["predicted_category"],
37         "cluster_id": result["cluster_id"],
38         "distance_to_centroid": result["distance_to_centroid"],
39         "distances_all_clusters": result["distances_all_clusters"],
40     })
41
42     return jsonify({"prediction_id": prediction_id, **result})
43
44
45 @api_bp.get("/dataset/stats")
46 def dataset_stats():
47     stats = []
48     for category in settings.CATEGORIES:
49         count = documents_col().count_documents({"category": category})
50         stats.append({
51             "category": category,
52             "required": settings.MIN_DOCS_PER_CATEGORY,
53             "actual": count,
54             "status": "PASS" if count >= settings.MIN_DOCS_PER_CATEGORY else "FAIL",
55         })
```

Source Code: task2_routes_api_part1

```

56     total = sum(s["actual"] for s in stats)
57     return jsonify({
58         "categories": stats,
59         "total": total,
60         "total_required": settings.MIN_DOCS_PER_CATEGORY * len(settings.CATEGORIES),
61         "total_status": "PASS" if total >= settings.MIN_DOCS_PER_CATEGORY * len(settings.CATEGORIES) else "FAIL",
62     })
63
64
65 @api_bp.get("/model/evaluation")
66 def model_evaluation():
67     if not settings.EVALUATION_REPORT_PATH.exists():
68         return jsonify({"error": "Model has not been trained yet."}), 503
69     with open(settings.EVALUATION_REPORT_PATH, "r", encoding="utf-8") as f:
70         return jsonify(json.load(f))
71
72 @api_bp.get("/model/pca")
73 def model_pca():
74     pca_path = settings.MODELS_DIR / "pca_data.json"
75     if not pca_path.exists():
76         return jsonify({"error": "PCA data not found. Run train_model.py."}), 404
77     with open(pca_path, "r", encoding="utf-8") as f:
78         return jsonify(json.load(f))
79
80
81 @api_bp.get("/predictions/history")
82 def predictions_history():
83     try:
84         limit = min(50, max(1, int(request.args.get("limit", 20))))
85     except ValueError:
86         limit = 20
87     docs = list(predictions_col().find({}).sort("timestamp", -1).limit(limit))
88     for d in docs:
89         d["_id"] = str(d["_id"])
90         ts = d["timestamp"]
91         if ts.tzinfo is None:
92             from datetime import timezone
93             ts = ts.replace(tzinfo=timezone.utc)
94         d["timestamp"] = ts.isoformat()
95     return jsonify({"predictions": docs})
96
97
98 @api_bp.delete("/predictions/history")
99 def clear_predictions_history():
100     predictions_col().delete_many({})
101     return jsonify({"message": "History cleared successfully."})
102
103 import re
104
105 @api_bp.get("/suggest")
106 def suggest():
107     query = request.args.get("q", "").strip()
108     if not query or len(query) < 2:
109         return jsonify([])
110

```

Source Code: task2_routes_api_part2

```

111     regex = re.compile(f"{re.escape(query)}", re.IGNORECASE)
112     suggestions = []
113
114     for d in documents_col().find({"content": regex}).limit(8):
115         if "content" in d:
116             text = d["content"]
117             snippet = text[:90] + ("..." if len(text) > 90 else "")
118             snippet = snippet.replace(&#x27;\n', ' ').strip()
119             suggestions.append(snippet)
120
121     return jsonify(list(dict.fromkeys(suggestions)))
122
123 @api_bp.post("/retrain")
124 def retrain_model():
125     import subprocess
126     import sys
127     import threading
128     from config.settings import settings
129
130     def run_scripts():
131         try:
132             logger.info("Starting dataset collection...")
133             subprocess.run([sys.executable, str(settings.BACKEND_DIR.parent / "scripts" / "build_dataset.py")], check=True)
134             logger.info("Starting model retraining...")
135             subprocess.run([sys.executable, str(settings.BACKEND_DIR.parent / "scripts" / "train_model.py")], check=True)
136             logger.info("Retraining complete.")
137         except Exception as e:
138             logger.error(f"Retrain pipeline failed: {e}")
139
140     thread = threading.Thread(target=run_scripts)
141     thread.daemon = True
142     thread.start()
143
144     return jsonify({"message": "Retraining pipeline started in the background."})
145

```

Source Code: task2_routes_api_part3

```

1  """Entry point for the Task 2 document clustering backend."""
2  import os
3  from app import create_app
4  from config.settings import settings
5
6  if __name__ == "__main__":
7      app = create_app()
8      port = int(os.environ.get("PORT", settings.TASK2_PORT))
9      app.run(host="0.0.0.0", port=port, debug=False, use_reloader=False)
10

```

Source Code: task2_run

```
build_dataset.py (Part 1)

1  """
2  Builds the Task 2 labelled corpus (Economics / Entertainment / Politics)
3  from the Greene & Cunningham BBC News full-text dataset.
4
5  Data source and licence (Option B: "a legitimate existing dataset" --
6  chosen deliberately over live scraping of bbc.co.uk, whose own robots.txt
7  explicitly states "No scraping, crawling, or systematic extraction of
8  content" and "No creating datasets from BBC content", and which blocks
9  several AI-related crawlers by name):
10
11      D. Greene and P. Cunningham. "Practical Solutions to the Problem of
12      Diagonal Dominance in Kernel Document Clustering." Proceedings of the
13      23rd International Conference on Machine Learning (ICML), 2006.
14      Dataset page: http://mlg.ucd.ie/datasets/bbc.html
15      "These datasets are made available for non-commercial and research
16      purposes only." Original articles (2004-2005): copyright BBC.
17
18  The dataset's "business" category is used for the coursework's
19  "Economics" category (news outlets commonly file macroeconomic/monetary
20  policy reporting under a "business" desk; the article content itself is
21  economics reporting).
22
23  For each of the three required categories, the TARGET_DOCS_PER_CATEGORY
24  longest available articles are kept (coursework guidance: "the longer is
25  usually the better"), each stored with its original filename as
26  document_id, its first line as title, full text as content, and the
27  academic citation above as source/source_url.
28  """
29  import csv
30  import json
31  import sys
32  from datetime import date, datetime, timezone
33  from pathlib import Path
34
35  sys.path.insert(0, str(Path(__file__).resolve().parents[1] / "backend"))
36
37  from config.settings import settings # noqa: E402
38  from database.mongo_client import documents_col, ensure_indexes # noqa: E402
39
40  CATEGORY_FOLDER_MAP = {
41      "Economics": "business",
42      "Entertainment": "entertainment",
43      "Politics": "politics",
44  }
45  SOURCE_CITATION = (
46      "Greene, D., & Cunningham, P. (2006). Practical Solutions to the "
47      "Problem of Diagonal Dominance in Kernel Document Clustering. "
48      "Proceedings of the 23rd International Conference on Machine "
49      "Learning (ICML 2006). "
50  )
51  SOURCE_URL = "http://mlg.ucd.ie/datasets/bbc.html"
52
53
54  def load_category_files(category: str) -> list[dict]:
55      folder = settings.RAW_DATASET_DIR / CATEGORY_FOLDER_MAP[category]
```

Source Code: task2_scripts_build_dataset_part1


```
build_dataset.py (Part 2)

56 docs = []
57 for path in sorted(folder.glob("*.txt")):
58     # The original archive is Windows-1252 encoded (curly quotes, £ signs).
59     raw = path.read_text(encoding="cp1252", errors="replace")
60     lines = raw.splitlines()
61     title = lines[0].strip() if lines else path.stem
62     content = raw.strip()
63     docs.append({
64         "document_id": f"{CATEGORY_FOLDER_MAP[category]}-{path.stem}",
65         "title": title,
66         "content": content,
67         "category": category,
68         "source": SOURCE_CITATION,
69         "source_url": SOURCE_URL,
70         "collection_date": datetime.now(timezone.utc).isoformat(),
71         "original_publication_period": "2004-2005",
72         "word_count": len(content.split()),
73     })
74 return docs
75
76
77 def select_top_n(docs: list[dict], n: int) -> list[dict]:
78     """Prefer longer, information-richer documents, as the brief advises."""
79     return sorted(docs, key=lambda d: d["word_count"], reverse=True)[:n]
80
81
82 def main():
83     ensure_indexes()
84
85     all_selected = []
86     report_rows = []
87
88     for category in settings.CATEGORIES:
89         available = load_category_files(category)
90         selected = select_top_n(available, settings.TARGET_DOCS_PER_CATEGORY)
91
92         all_selected.extend(selected)
93
94         status = "PASS" if len(selected) >= settings.MIN_DOCS_PER_CATEGORY else "FAIL"
95         report_rows.append({
96             "category": category,
97             "required": settings.MIN_DOCS_PER_CATEGORY,
98             "available_in_source": len(available),
99             "actual_selected": len(selected),
100             "status": status,
101         })
102
103     total = len(all_selected)
104     total_status = "PASS" if total >= settings.MIN_DOCS_PER_CATEGORY * len(settings.CATEGORIES) else "FAIL"
105     report_rows.append({
106         "category": "TOTAL",
107         "required": settings.MIN_DOCS_PER_CATEGORY * len(settings.CATEGORIES),
108         "available_in_source": sum(r["available_in_source"] for r in report_rows),
109         "actual_selected": total,
110         "status": total_status,
```

Source Code: [task2_scripts_build_dataset_part2](#)

```
build_dataset.py (Part 3)

111     })
112
113     print("\nDataset validation report")
114     print(f"{'&#x27;Category':<15}{'Required':>10}{'Available':>12}{'Selected':>10}{'Status':>8}")
115     for row in report_rows:
116         print(f"{'row[&#x27;category']':<15}{'row['required']':>10}{'row['available_in_source']':>12}{'row['actual_select
ed']':>10}{'row['status']':>8}")
117
118     if any(r["status"] == "FAIL" for r in report_rows):
119         print("\nDataset validation FAILED. Aborting before writing to MongoDB.")
120         sys.exit(1)
121
122     # Persist to MongoDB (upsert by document_id so re-running is idempotent).
123     inserted, updated = 0, 0
124     all_doc_ids = []
125     for doc in all_selected:
126         all_doc_ids.append(doc["document_id"])
127         result = documents_col().update_one(
128             {"document_id": doc["document_id"]}, {"$set": doc}, upsert=True
129         )
130         if result.upserted_id is not None:
131             inserted += 1
132         else:
133             updated += 1
134
135     # Delete documents that are no longer in the dataset (e.g., cleared history items)
136     delete_result = documents_col().delete_many({"document_id": {"$nin": all_doc_ids}})
137     deleted = delete_result.deleted_count
138
139     # Also write a local CSV/JSON snapshot for the processed dataset folder,
140     # and a machine-readable validation report used by tests / documentation.
141     csv_path = settings.PROCESSED_DATASET_DIR / "dataset.csv"
142     with open(csv_path, "w", newline="", encoding="utf-8") as f:
143         writer = csv.DictWriter(f, fieldnames=list(all_selected[0].keys()))
144         writer.writeheader()
145         writer.writerows(all_selected)
146
147     report_path = settings.PROCESSED_DATASET_DIR / "dataset_validation_report.json"
148     with open(report_path, "w", encoding="utf-8") as f:
149         json.dump({"generated_at": datetime.now(timezone.utc).isoformat(), "rows": report_rows}, f, indent=2)
150
151     print(f"\nMongoDB: {inserted} inserted, {updated} updated in &#x27;clustering_documents'")
152     print(f"CSV snapshot written to {csv_path}")
153     print(f"Validation report written to {report_path}")
154
155
156 if __name__ == "__main__":
157     main()
158
```

Source Code: task2_scripts_build_dataset_part3

```
train_model.py

1  """
2  Trains the TF-IDF + K-Means clustering model on the corpus stored in
3  MongoDB, evaluates it against the known category labels, generates the
4  2D cluster visualisation, and persists the vectorizer/model/mapping to
5  disk so the running API never retrains on a live request.
6  """
7  import json
8  import sys
9  from pathlib import Path
10
11  sys.path.insert(0, str(Path(__file__).resolve().parents[1] / "backend"))
12
13  from clustering.kmeans_model import train # noqa: E402
14  from config.settings import settings # noqa: E402
15  from database.mongo_client import documents_col # noqa: E402
16  from utils.logging_setup import configure_logging # noqa: E402
17  from visualization.pca_plot import generate_cluster_plot # noqa: E402
18
19
20  def main():
21      configure_logging()
22
23      documents = list(documents_col().find({}, {"content": 1, "category": 1, "title": 1, "document_id": 1}))
24      if not documents:
25          print("No documents found in MongoDB. Run scripts/build_dataset.py first.")
26          sys.exit(1)
27
28      print(f"Loaded {len(documents)} documents from MongoDB for training.")
29      result = train(documents)
30
31      # Pass the original texts and true labels to the visualisation
32      plot_path = generate_cluster_plot(
33          X=result["X"],
34          cluster_labels=result["cluster_labels"],
35          true_categories=result["true_categories"],
36          cluster_to_category=result["cluster_to_category"],
37          titles=[d.get("title", "") for d in result["documents"]]
38      )
39
40      print("\n=== Cluster -> Category mapping ===")
41      print(json.dumps(result["cluster_to_category"], indent=2))
42
43      print("\n=== Evaluation ===")
44      print(json.dumps(result["evaluation"], indent=2))
45
46      print(f"\nVisualisation saved to: {plot_path}")
47      print(f"Model artifacts saved to: {settings.MODELS_DIR}")
48
49
50  if __name__ == "__main__":
51      main()
52
```

Source Code: task2_scripts_train_model

```

1  """
2  2D visualisation of the K-Means clustering result.
3
4  TF-IDF vectors are extremely high-dimensional (one dimension per
5  vocabulary term), so they cannot be plotted directly. Principal Component
6  Analysis (PCA) is used purely to project each document vector down to the
7  2 directions of greatest variance for plotting -- PCA plays no role in
8  the clustering decision itself (K-Means was already fit on the full
9  TF-IDF matrix in clustering/kmeans_model.py).
10 """
11 import logging
12
13 import matplotlib
14 matplotlib.use("Agg") # headless rendering, no display required
15 import matplotlib.pyplot as plt
16 from sklearn.decomposition import PCA
17
18 from config.settings import settings
19
20 logger = logging.getLogger("task2.visualization")
21
22 CATEGORY_COLORS = {"Economics": "#3b82f6", "Entertainment": "#a855f7", "Politics": "#10b981"}
23
24 import json
25 def generate_cluster_plot(X, cluster_labels, true_categories, cluster_to_category, titles=None, out_path=None) -> s
    tr:
26     out_path = out_path or (settings.FIGURES_DIR / "figure_task2_kmeans_clusters.png")
27
28     pca = PCA(n_components=2, random_state=42)
29     coords = pca.fit_transform(X.toarray())
30
31     fig, ax = plt.subplots(figsize=(9, 5), facecolor=&#x27;none')
32     ax.set_facecolor(&#x27;none')
33     ax.grid(color=&#x27;#2a2d3e', linestyle='-', linewidth=0.5, zorder=0)
34
35     for category in settings.CATEGORIES:
36         mask = [c == category for c in true_categories]
37         xs = [coords[i, 0] for i, m in enumerate(mask) if m]
38         ys = [coords[i, 1] for i, m in enumerate(mask) if m]
39         ax.scatter(xs, ys, label=f"{category}", color=CATEGORY_COLORS[category],
40                   s=35, alpha=0.9, edgecolor="none", zorder=3)
41
42     explained = pca.explained_variance_ratio_
43     ax.set_xlabel(f"PCA Component 1", color=&#x27;#64748b', fontsize=9)
44     ax.set_ylabel(f"PCA Component 2", color=&#x27;#64748b', fontsize=9)
45     ax.tick_params(colors=&#x27;none') # Hide ticks
46
47     for spine in ax.spines.values():
48         spine.set_color(&#x27;#2a2d3e')
49
50     fig.tight_layout()
51     fig.savefig(out_path, dpi=150, transparent=True)
52     plt.close(fig)
53
54     # Save PCA coordinates to JSON for frontend interactive plotting
55     pca_data = []

```

Source Code: [task2_visualization_pca_part1](#)

Appendix D: Pytest Evidence

```
PowerShell — pytest -v (task1_vertical_search/backend)

===== test session starts =====
platform win32 -- Python 3.14.7, pytest-9.1.1, pluggy-1.6.0 -- C:\Users\jrohi\OneDrive\Desktop\My
INFORETRIVAL_ASSIGNMENT_Claude\IR_COURSEWORK\.venv\Scripts\python.exe
cachedir: .pytest_cache
rootdir: C:\Users\jrohi\OneDrive\Desktop\My INFORETRIVAL_ASSIGNMENT_Claude\IR_COURSEWORK\task1_vertical_search\backend
collecting ... collected 29 items

tests/test_api.py::test_search_requires_query_param PASSED [ 3%]
tests/test_api.py::test_search_returns_top_k_page_size PASSED [ 6%]
tests/test_api.py::test_search_nonsense_query_returns_empty_results_not_error PASSED [ 10%]
tests/test_api.py::test_search_result_shape_includes_required_fields PASSED [ 13%]
tests/test_api.py::test_research_output_invalid_id_returns_400 PASSED [ 17%]
tests/test_api.py::test_research_output_missing_id_returns_404 PASSED [ 20%]
tests/test_api.py::test_crawler_status_endpoint_shape PASSED [ 24%]
tests/test_database.py::test_upsert_by_url_prevents_duplicates PASSED [ 27%]
tests/test_database.py::test_mongodb_connection_is_reachable PASSED [ 31%]
tests/test_parsers.py::test_extract_links_finds_publications_and_persons PASSED [ 34%]
tests/test_parsers.py::test_parse_publication_detail_extracts_required_fields PASSED [ 37%]
tests/test_parsers.py::test_parse_profile_detail_extracts_required_fields PASSED [ 41%]
tests/test_preprocessing.py::test_lowercases_and_removes_punctuation PASSED [ 44%]
tests/test_preprocessing.py::test_removes_stopwords PASSED [ 48%]
tests/test_preprocessing.py::test_empty_input_returns_empty_string PASSED [ 51%]
tests/test_preprocessing.py::test_same_pipeline_applies_to_query_and_document_text PASSED [ 55%]
tests/test_robots.py::test_seed_organisation_page_is_allowed PASSED [ 58%]
tests/test_robots.py::test_publication_and_person_pages_are_allowed PASSED [ 62%]
tests/test_robots.py::test_rss_export_query_formats_are_disallowed PASSED [ 65%]
tests/test_robots.py::test_crawl_delay_is_five_seconds PASSED [ 68%]
tests/test_scheduler.py::test_scheduler_job_interval_matches_configured_weeks PASSED [ 72%]
tests/test_scheduler.py::test_scheduler_status_reports_running_state PASSED [ 75%]
tests/test_vector_space_model.py::test_index_builds_from_available_documents PASSED [ 79%]
tests/test_vector_space_model.py::test_keyword_search_ranks_relevant_document_first PASSED [ 82%]
tests/test_vector_space_model.py::test_author_name_search_returns_matching_publication PASSED [ 86%]
tests/test_vector_space_model.py::test_results_sorted_descending_by_cosine_similarity PASSED [ 89%]
```

Figure 25 Full Pytest execution for task 1

```
PowerShell — pytest -v (task2_document_clustering/backend)

===== test session starts =====
platform win32 -- Python 3.14.7, pytest-9.1.1, pluggy-1.6.0 -- C:\Users\jrohi\OneDrive\Desktop\My
INFORETRIVAL_ASSIGNMENT_Claude\IR_COURSEWORK\.venv\Scripts\python.exe
cachedir: .pytest_cache
rootdir: C:\Users\jrohi\OneDrive\Desktop\My INFORETRIVAL_ASSIGNMENT_Claude\IR_COURSEWORK\task2_document_clustering\backend
collecting ... collected 21 items


tests/test_api.py::test_classify_economics_example_from_brief PASSED [ 4%]
tests/test_api.py::test_classify_entertainment_example PASSED [ 9%]
tests/test_api.py::test_classify_empty_input_returns_400 PASSED [ 14%]
tests/test_api.py::test_classify_missing_field_returns_400 PASSED [ 19%]
tests/test_api.py::test_classify_invalid_json_body_returns_400 PASSED [ 23%]
tests/test_api.py::test_classify_long_input_is_accepted PASSED [ 28%]
tests/test_api.py::test_classify_over_max_length_returns_400 PASSED [ 33%]
tests/test_api.py::test_dataset_stats_endpoint PASSED [ 38%]
tests/test_api.py::test_model_evaluation_endpoint PASSED [ 42%]
tests/test_api.py::test_predictions_history_endpoint PASSED [ 47%]
tests/test_dataset_validation.py::test_validation_report_file_exists_and_passes PASSED [ 52%]
tests/test_dataset_validation.py::test_mongodb_has_at_least_minimum_documents_per_category PASSED [ 57%]
tests/test_dataset_validation.py::test_total_documents_meets_overall_minimum PASSED [ 61%]
tests/test_kmeans_model.py::test_train_produces_three_clusters_and_full_mapping PASSED [ 66%]
tests/test_kmeans_model.py::test_train_evaluation_reports_required_metrics PASSED [ 71%]
tests/test_kmeans_model.py::test_classify_uses_persisted_model_without_retraining PASSED [ 76%]
tests/test_kmeans_model.py::test_empty_and_whitespace_text_still_produce_a_prediction PASSED [ 80%]
tests/test_preprocessing.py::test_pipeline_lowercases_removes_punctuation_and_stopwords PASSED [ 85%]
tests/test_preprocessing.py::test_stemming_collapses_related_word_forms PASSED [ 90%]
tests/test_preprocessing.py::test_empty_input_returns_empty_string PASSED [ 95%]
tests/test_preprocessing.py::test_long_input_is_processed_without_error PASSED [100%]

===== warnings summary =====
tests/test_api.py::test_classify_economics_example_from_brief
tests/test_api.py::test_classify_economics_example_from_brief
tests/test_api.py::test_classify_economics_example_from_brief
tests/test_kmeans_model.py::test_classify_uses_persisted_model_without_retraining
tests/test_kmeans_model.py::test_classify_uses_persisted_model_without_retraining
tests/test_kmeans_model.py::test_classify_uses_persisted_model_without_retraining
tests/test_kmeans_model.py::test_empty_and_whitespace_text_still_produce_a_prediction
tests/test_kmeans_model.py::test_empty_and_whitespace_text_still_produce_a_prediction
tests/test_kmeans_model.py::test_empty_and_whitespace_text_still_produce_a_prediction
C:\Users\jrohi\OneDrive\Desktop\My INFORETRIVAL_ASSIGNMENT_Claude\IR_COURSEWORK\.venv\Lib\site-packages\joblib\numpy_pickle.py:207: DeprecationWarning:
Setting the shape on a NumPy array has been deprecated in NumPy 2.5.
```


Figure 26 Full Pytest execution for Task 2.

Appendix E: Additional Application Screenshots

The following screenshots represent additional UI views, visualizations, and MongoDB states captured during the system's live operation.




Additional System Screenshot: 01_unified_search_home




Additional System Screenshot: 02_unified_search_results

Task 1 — Search Results: "digital health" (39 results)		
#1	Rethinking Digital Healthcare Education for Frailty Authors: Jo Jennings, Juliana Samson Published: 2025	Cosine sim: 0.2330
#2	A digital yoga-based 'Pause Grow Reflect' intervention for postgraduate researchers: a feasibil... Authors: Natalie Bisal, Riya Patel, Nikki Holliday... Published: 12 Jun 2026	Cosine sim: 0.1161
#3	A novel digital intervention to support mental and sexual wellbeing after stroke and brain inju... Authors: Hayley Wright, Andy Turner, David McWilliams... Published: 21 Apr 2026	Cosine sim: 0.1064
#4	Five Insights from the Global Burden of Disease Study 2019 Authors: GBD 2019 Viewpoint Collaborators Published: 17 Oct 2020	Cosine sim: 0.0932
#5	A verbal and visual account of the perceived influences on health-promoting behaviors among wom... Authors: Maxine E Whelan, Mark W Orme, Karen Pearce... Published: 7 Apr 2026	Cosine sim: 0.0519


Additional System Screenshot: 03_task1_search_digital_health



Additional System Screenshot: 03_unified_search_pagination




Additional System Screenshot: 04_task2_dataset_stats




Additional System Screenshot: 04_unified_crawler_status

Task 2 — K-Means Model Evaluation Report	
Documents / vocabulary terms	540 / 290
Number of clusters (K)	3
Accuracy (vs known labels)	93.0%
Precision (macro avg)	0.9333
Recall (macro avg)	0.9296
F1 Score (macro avg)	0.9296
Silhouette Score	0.0604
Inertia	431.15


Additional System Screenshot: 05_task2_model_evaluation




Additional System Screenshot: 05_unified_cluster_home




Additional System Screenshot: 06_unified_cluster_result



Additional System Screenshot: 08_unified_model_evaluation



Additional System Screenshot: 08_unified_search_suggestions



Additional System Screenshot: new_02_crawler_status